

# RDF Dataset Canonicalization

## A Standard RDF Dataset Canonicalization Algorithm

W3C Recommendation 21 May 2024



### ▼ More details about this document

#### This version:

<https://www.w3.org/TR/2024/REC-rdf-canon-20240521/>

#### Latest published version:

<https://www.w3.org/TR/rdf-canon/>

#### Latest editor's draft:

<https://w3c.github.io/rdf-canon/spec/>

#### History:

<https://www.w3.org/standards/history/rdf-canon/>

[Commit history](#)

#### Test suite:

<https://w3c.github.io/rdf-canon/tests/>

#### Implementation report:

<https://w3c.github.io/rdf-canon/reports/>

#### Editors:

[Dave Longley \(Digital Bazaar\)](#)

[Gregg Kellogg](#)

[Dan Yamamoto](#)

#### Former editor:

[Manu Sporny \(Digital Bazaar\) \(CG Report\)](#)

#### Author:

[Dave Longley \(Digital Bazaar\)](#)

#### Feedback:

[GitHub w3c/rdf-canon \(pull requests, new issue, open issues\)](#)

[public-rch-wg@w3.org](mailto:public-rch-wg@w3.org) with subject line [rdf-canon] ... message topic ... ([archives](#))

#### Errata:

[Errata exists.](#)

See also [translations](#).

[Copyright](#) © 2010-2024 [World Wide Web Consortium](#). W3C® [liability](#), [trademark](#) and [permissive document license](#) rules apply.

---

## Abstract

RDF [RDF11-CONCEPTS] describes a graph-based data model for making claims about the world and provides the foundation for reasoning upon that graph of information. At times, it becomes necessary to compare the differences between sets of graphs, digitally sign them, or generate short identifiers for graphs via hashing algorithms. This document outlines an algorithm for normalizing [RDF datasets](#) such that these operations can be performed.

## Status of This Document

*This section describes the status of this document at the time of its publication. A list of current W3C publications and the latest revision of this technical report can be found in the [W3C technical reports index](#) at <https://www.w3.org/TR/>.*

This document describes the [RDFC-1.0](#) algorithm for canonicalizing RDF datasets, which was the input from the [W3C Credentials Community Group](#) published as [\[CCG-RDC-FINAL\]](#).

At the time of publication, [RDF11-CONCEPTS] is the most recent recommendation defining [RDF datasets](#) and [\[N-QUADS\]](#), however work on an updated specification is ongoing within the [W3C RDF-star Working Group](#). Some dependencies from relevant updated specifications are provided normatively in this specification with the expectation that a future update to this specification will replace those with normative references to updated RDF specifications.

This document was published by the [RDF Dataset Canonicalization and Hash Working Group](#) as a Recommendation using the [Recommendation track](#).

W3C recommends the wide deployment of this specification as a standard for the Web.

A W3C Recommendation is a specification that, after extensive consensus-building, is endorsed by W3C and its Members, and has commitments from Working Group members to [royalty-free licensing](#) for implementations. Future updates to this Recommendation may incorporate [new features](#).

This document was produced by a group operating under the [W3C Patent Policy](#). W3C maintains a [public list of any patent disclosures](#) made in connection with the deliverables of the group; that page also includes instructions for disclosing a patent. An individual who has actual knowledge of a patent

which the individual believes contains [Essential Claim\(s\)](#) must disclose the information in accordance with [section 6 of the W3C Patent Policy](#).

This document is governed by the [03 November 2023 W3C Process Document](#).

## Table of Contents

### **Abstract**

### **Status of This Document**

## **1. Introduction**

- 1.1      Uses of Dataset Canonicalization
- 1.2      How to Read this Document
- 1.3      Typographical conventions

## **2. Conformance**

## **3. Terminology**

- 3.1      Terms defined by this specification
- 3.2      Terms defined by cited specifications

## **4. Canonicalization**

- 4.1      Overview
- 4.2      Canonicalization State
- 4.3      Blank Node Identifier Issuer State
- 4.4      Canonicalization Algorithm
  - 4.4.1      Overview
  - 4.4.2      Examples
  - 4.4.3      Algorithm
- 4.5      Issue Identifier Algorithm
  - 4.5.1      Overview
  - 4.5.2      Algorithm
- 4.6      Hash First Degree Quads
  - 4.6.1      Overview
  - 4.6.2      Examples
  - 4.6.3      Algorithm
- 4.7      Hash Related Blank Node

- 4.7.1 Overview
- 4.7.2 Examples
- 4.7.3 Algorithm
- 4.8 Hash N-Degree Quads
  - 4.8.1 Overview
  - 4.8.2 Examples
  - 4.8.3 Algorithm

## 5. **Serialization**

## 6. **Privacy Considerations**

## 7. **Security Considerations**

- 7.1 Dataset Poisoning
- 7.2 Insecure Hash Algorithms

## 8. **Use Cases**

## 9. **Examples**

- 9.1 Duplicate Paths
- 9.2 Double Circle
- 9.3 Dataset with Blank Node Named Graph

## A. **A Canonical form of N-Quads**

## B. **URDNA2015**

## C. **URGNA2012**

## D. **Index**

- D.1 Terms defined by this specification
- D.2 Terms defined by reference

## E. **Changes since the First Public Working Draft of 24 November 2022**

## F. **Changes since the Candidate Recommendation Snapshot of 31 October 2023**

## G. **Acknowledgements**

## H. **References**

- H.1 Normative references
- H.2 Informative references

## § 1. Introduction

*This section is non-normative.*

When data scientists discuss canonicalization, they do so in the context of achieving a particular set of goals. Since the same information may sometimes be expressed in a variety of different ways, it often becomes necessary to transform each of these different ways into a single, standard representation. With a standard representation, the differences between two different sets of data can be easily determined, a cryptographically-strong hash identifier can be generated for a particular set of data, and a particular set of data may be digitally-signed for later verification.

In particular, this specification is about normalizing [RDF datasets](#), which are collections of graphs. Since a directed graph can express the same information in more than one way, it requires canonicalization to achieve the aforementioned goals and any others that may arise via serendipity.

Most [RDF datasets](#) can be canonicalized fairly quickly, in terms of algorithmic time complexity. However, those that contain nodes that do not have globally unique identifiers pose a greater challenge. Normalizing these datasets presents the ***graph isomorphism*** problem, a problem that is believed to be difficult to solve quickly in the worst case. Fortunately, existing real world data is rarely, if ever, modeled in a way that manifests as the worst case and new data can be modeled to avoid it. In fact, software systems that detect a problematic dataset (see [7.1 Dataset Poisoning](#)) can choose to assume it's an attempted denial of service attack, rather than a real input, and abort.

This document outlines an algorithm for generating a canonical serialization of an [RDF dataset](#) given an [RDF dataset](#) as input. The algorithm is called the **RDF Canonicalization algorithm version 1.0** or [RDPC-1.0](#).

## NOTE

[RDF 1.1 Concepts and Abstract Syntax \[RDF11-CONCEPTS\]](#) lacks clarity on the representation of [language-tagged strings](#), where [language tags](#) of the form `xx-YY` are treated as being case insensitive. Implementations might represent language tags using all lower case in the form `xx-yy`, retain the original representation `xx-YY`, or use [BCP47] [formatting conventions](#), leading to different canonical forms, and therefore, different hashed values.

- The Canonicalization algorithm is based on the RDF 1.1 definition, in the sense that the [language tag](#) `xx-YY` is case insensitive, which might lead to different canonicalizations if the user is not aware of this problem.
- User communities ought to agree to use lower case [language tags](#), while being aware that some implementations might normalize language tags, affecting hash values.
- Future evolution of RDF might regulate this issue, which RDF environments might have to adapt to, and this might lead to an update of [RDPC-1.0](#).

## NOTE

See [B. URDNA2015](#) for a comparison with the version of the algorithm published in [RDF Dataset Canonicalization \[CCG-RDC-FINAL\]](#).

## § 1.1 Uses of Dataset Canonicalization

There are different use cases where graph or dataset canonicalization are important:

- Determining if one serialization is isomorphic to another.
- Digital signing of graphs (datasets) independent of serialization or format.
- Comparing two graphs (datasets) to find differences.
- Communicating change sets when remotely updating an [RDF source](#).

A canonicalization algorithm is necessary, but not necessarily sufficient, to handle many of these use cases. The use of [blank nodes](#) in RDF graphs and datasets has a long history and creates inevitable complexities. Blank nodes are used for different purposes:

- when a well known identifier for a node is not known, or the author of a document chooses not to unambiguously name that node,
- when a node is used to stitch together parts of a graph and the nodes themselves are not interesting (e.g., [RDF Collections](#) in [[RDF11-MT](#)]),
- when someone is trying to create an intentionally difficult graph topology.

Furthermore, RDF semantics dictate that deserializing an RDF document results in the creation of unique [blank nodes](#), unless it can be determined that on each occasion, the [blank node](#) identifies the same resource. This is due to the fact that [blank node identifiers](#) are an aspect of a concrete RDF syntax and are not intended to be persistent or portable. Within the abstract RDF model, blank nodes do not have identifiers (although some RDF store implementations may use stable identifiers and may choose to make them portable). See [Blank Nodes](#) in [[RDF11-CONCEPTS](#)] for more information.

RDF does have a provision for allowing blank nodes to be published in an externally identifiable way through the use of [Skolem IRIs](#), which allow a given RDF store to replace the use of blank nodes in a concrete syntax with [IRIs](#), which then serve to repeatably identify that blank node within that particular RDF store; however, this is not generally useful for talking about the same graph in different RDF stores, or other concrete representations. In any case, a stable [blank node identifier](#) defined for one RDF store or serialization is arbitrary, and typically not relatable to the context within which it is used.

This specification defines an algorithm for creating stable [blank node identifiers](#) repeatably for different serializations possibly using individualized [blank node identifiers](#) of the same RDF graph (dataset) by grounding each [blank node](#) through the nodes to which it is connected. As a result, a graph signature can be obtained by hashing a canonical serialization of the resulting [canonicalized dataset](#), allowing for the isomorphism and digital signing use cases. This specification does not define such a graph signature.

As blank node identifiers can be stable even with other changes to a graph (dataset), in some cases it is possible to compute the difference between two graphs (datasets), for example if changes are made only to ground triples, or if new blank nodes are introduced which do not create an automorphic confusion with other existing blank nodes. If any information which would change the generated blank node identifier, a resulting diff might indicate a greater set of changes than actually exists. Additionally, if the starting dataset is an N-Quads document, it may be possible to correlate the original blank node identifiers used within that N-Quads document with those issued in the [canonicalized dataset](#).

**NOTE**

Although alternative [hash algorithms](#) might be used with this specification, applications ought to carefully weigh the advantages and disadvantages of using an alternative hash function. This is the case, in particular, for any representation of the [canonical n-quads form](#) or [issued identifiers map](#) that does not identify the associated hash algorithm. Any use case that requires reproduction of the same output is expected to unequivocally express or communicate the internal hash algorithm that was used when generating the [canonical n-quads form](#).

## § 1.2 How to Read this Document

This document is a detailed specification for an [RDF dataset](#) canonicalization algorithm. The document is primarily intended for the following audiences:

- Software developers that want to implement an [RDF dataset](#) canonicalization algorithm.
- Masochists.

To understand the basics in this specification you must be familiar with basic RDF concepts [[RDF11-CONCEPTS](#)]. A working knowledge of [graph theory](#) and [graph isomorphism](#) is also recommended.

## § 1.3 Typographical conventions

*This section is non-normative.*

The following typographic conventions are used in this specification:

**markup**

Markup (elements, attributes, properties), machine processable values (string, characters, media types), property names, and file names are in red-orange monospace font.

**variable**

A variable in pseudo-code or in an algorithm description is italicized.

**definition**

A definition of a term, to be used elsewhere in this or other specifications, is italicized and in bold.

## **definition reference**

A reference to a definition *in this document* is underlined and is also an active link to the definition itself.

## **markup definition reference**

References to a definition *in this document*, when the reference itself is also a markup, is underlined, in a red-orange monospace font, and is also an active link to the definition itself.

## **external definition reference**

A reference to a definition *in another document* is underlined and italicized, and is also an active link to the definition itself.

## **markup external definition reference**

A reference to a definition *in another document*, when the reference itself is also a markup, is underlined and italicized in a red-orange monospace font, and is also an active link to the definition itself.

## **hyperlink**

A hyperlink is underlined and in blue.

## **[reference]**

A document reference (normative or informative) is enclosed in square brackets and links to the references section.

### **Explanation**

An expandable area to find a more detailed, non-normative explanation of a particular algorithmic step.

#### ► Explanation

### **Logging**

An expandable area to find suggestions for implementations to log information about processing, which may be useful in comparing with other implementations, or with logs provided with each test case.

#### ► Logging

### **NOTE**

Notes are in light green boxes with a green left border and with a "Note" header in green. Notes are always informative.

### EXAMPLE 1

Examples are in light khaki boxes, with khaki left border, and with a numbered "Example" header in khaki.

Examples are always informative. The content of the example is in monospace font syntax colored.

Examples may have tabbed navigation buttons to show the results of transforming an example into other representations.

Code examples are generally given in a Turtle or TriG format for brevity, where each line represents a single triple or quad.

Additionally, have the following implied directives:

```
BASE <http://example.com/>
PREFIX : <#>
```

Following the Turtle/TriG syntax rules, blank nodes always appear in the `\_:xyz` format.

## § 2. Conformance

As well as sections marked as non-normative, all authoring guidelines, diagrams, examples, and notes in this specification are non-normative. Everything else in this specification is normative.

The key words *MUST*, *MUST NOT*, and *SHOULD* in this document are to be interpreted as described in [BCP 14](#) [RFC2119] [RFC8174] when, and only when, they appear in all capitals, as shown here.

A conforming processor is a system which can generate the [canonical n-quads form](#) of an [input dataset](#) consistent with the algorithms defined in this specification.

The algorithms in this specification are normative, because to consistently reproduce the same canonical identifiers, implementations *MUST* strictly conform to the steps outlined in these algorithms.

**NOTE**

Implementers can partially check their level of conformance with this specification by successfully passing the test cases of the [RDF Dataset Canonicalization test suite](#). Note, however, that passing all the tests in the test suite does not imply complete conformance to this specification. It only implies that the implementation conforms to the aspects tested by the test suite.

## § 3. Terminology

### § 3.1 Terms defined by this specification

#### *canonical n-quads form*

The canonicalized representation of a quad is defined in [A. A Canonical form of N-Quads](#). A quad in [canonical n-quads form](#) represents a [graph name](#), if present, in the same manner as a [subject](#), and each quad is terminated with a single **LF** (line feed, code point **U+000A**).

#### *canonicalization function*

A [canonicalization function](#) maps [RDF datasets](#) into [isomorphic datasets](#) [RDF11-CONCEPTS]. Two datasets produce the same canonical result if and only if they are [isomorphic](#). The [RDFC-1.0](#) algorithm implements a [canonicalization function](#). Some datasets may be constructed to prevent this algorithm from terminating in a reasonable amount of time (see [7.1 Dataset Poisoning](#)), in which case the algorithm can be considered to be a *partial canonicalization function*.

#### *canonicalized dataset*

A [canonicalized dataset](#) is the combination of the following:

- an [RDF dataset](#) — the [input dataset](#),
- the [input blank node identifier map](#) — mapping [blank nodes](#) in the input dataset to [blank node identifiers](#), and
- the [issued identifiers map](#) from the [canonical issuer](#) — mapping identifiers in the input dataset to canonical identifiers

A concrete serialization of a [canonicalized dataset](#) *MUST* label all [blank nodes](#) using the canonical [blank node identifiers](#).

#### *gossip path*

A particular enumeration of every incident [mention](#) emanating from a blank node. This recursively includes transitively related [mentions](#) until any named node or blank node already labeled by a particular [identifier issuer](#) is reached. Gossip paths are encoded and operated on in

the [RDFC-1.0](#) algorithm as strings. (See [4.8 Hash N-Degree Quads](#) for more information on the construction of gossip paths.)

### **hash**

The lowercase, hexadecimal representation of a message digest.

### **hash algorithm**

The default hash algorithm used by [RDFC-1.0](#), namely, SHA-256 [[FIPS-180-4](#)].

Implementations *MUST* support a parameter to define the [hash algorithm](#), *MUST* support SHA-256 and SHA-384 [[FIPS-180-4](#)], and *SHOULD* support the ability to specify other hash algorithms. Using a different hash algorithm will generally result in different output than using the default.

#### NOTE

There is no expectation that the default hash algorithm will also be used by any application creating a hash digest of the canonical N-Quads result.

### **identifier issuer**

An identifier issuer is used to issue new [blank node identifiers](#). It maintains a [blank node identifier issuer state](#).

### **input blank node identifier map**

Records any blank node identifiers already assigned to the [input dataset](#). If the [input dataset](#) is provided as an N-Quads document, the [map](#) relates blank nodes in the abstract [input dataset](#) to the blank node identifiers used within the N-Quads document, otherwise, identifiers are assigned arbitrarily for each blank node in the input dataset not previously identified.

#### NOTE

Implementations or environments might deal with blank node identifiers more directly; for example, some implementations might retain blank node identifiers in the parsed or abstract dataset. Implementations are expected to reuse these to enable usable mappings between input blank node identifiers and output blank node identifiers outside of the algorithm.

### **input dataset**

The abstract [RDF dataset](#) that is provided as input to the algorithm.

### **mention**

A node is mentioned in a [quad](#) if it is a component of that quad, as a [subject](#), [predicate](#), [object](#), or [graph name](#).

### **mention set**

The set of all [quads](#) in a [dataset](#) that [mention](#) a node  $n$  is called the [mention set](#) of  $n$ , denoted  $Q_n$ .

***quad***

A tuple composed of [subject](#), [predicate](#), [object](#), and [graph name](#). This is a generalization of an [RDF triple](#) along with a [graph name](#).

## § 3.2 Terms defined by cited specifications

***blank node***

A [blank node](#) as specified by [RDF11-CONCEPTS]. In short, it is a node in a graph that is neither an [IRI](#), nor a [literal](#).

***blank node identifier***

A [blank node identifier](#) as specified by [RDF11-CONCEPTS]. In short, it is a [string](#) that begins with `_:` that is used as an identifier for a [blank node](#). [Blank node identifiers](#) are typically implementation-specific local identifiers; this document specifies an algorithm for deterministically specifying them.

Concrete syntaxes, like [Turtle] or [N-Quads], prepend blank node identifiers with the `_:` string to differentiate them from other nodes in the graph. This affects the canonicalization algorithm, which is based on calculating a hash over the representations of [quads](#) in this format.

***default graph***

The [default graph](#) as specified by [RDF11-CONCEPTS].

***graph name***

A [graph name](#) as specified by [RDF11-CONCEPTS].

***IRI***

An [IRI](#) (Internationalized Resource Identifier) is a string that conforms to the syntax defined in [RFC3987].

***object***

An [object](#) as specified by [RDF11-CONCEPTS].

***predicate***

A [predicate](#) as specified by [RDF11-CONCEPTS].

***RDF dataset***

A [dataset](#) as specified by [RDF11-CONCEPTS]. For the purposes of this specification, an [RDF dataset](#) is considered to be a set of [quads](#)

***RDF graph***

An [RDF graph](#) as specified by [RDF11-CONCEPTS].

***RDF triple***

A [triple](#) as specified by [RDF11-CONCEPTS].

**string**

A string is a sequence of zero or more Unicode characters.

**subject**

A subject as specified by [RDF11-CONCEPTS].

**true and false**

Values that are used to express one of two possible boolean states.

**Unicode code point order**

This refers to determining the order of two Unicode strings (**A** and **B**), using Unicode Codepoint Collation, as defined in [XPATH-FUNCTIONS], which defines a total ordering of strings comparing code points. Note that for UTF-8 encoded strings, comparing the byte sequences gives the same result as code point order.

## § 4. Canonicalization

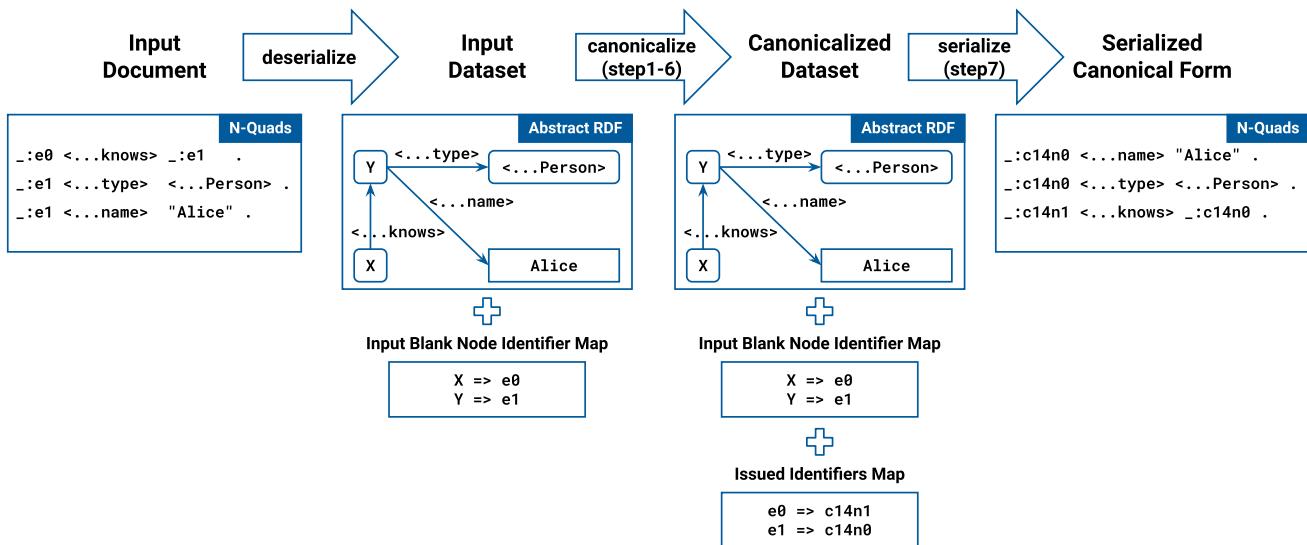
Canonicalization is the process of transforming an input dataset to its serialized canonical form. That is, any two input datasets that contain the same information, regardless of their arrangement, will be transformed into the same serialized canonical form. The problem requires directed graphs to be deterministically ordered into sets of nodes and edges. This is easy to do when all of the nodes have globally-unique identifiers, but can be difficult to do when some of the nodes do not. Any nodes without globally-unique identifiers must be issued deterministic identifiers.

### NOTE

This specification defines a canonicalized dataset to include stable identifiers for blank nodes, practical uses of which will always generate a canonical serialization of such a dataset.

In time, there may be more than one canonicalization algorithm and, therefore, for identification purposes, this algorithm is named the "RDF Canonicalization algorithm version 1.0" (**RDFC-1.0**).

Figure 1 provides an overview of RDFC-1.0, with steps 1 through 7 corresponding to the various steps described in 4.4.3 Algorithm.



*Figure 1* An illustrated overview of the [RDPC-1.0](#) algorithm.

*Image available in [SVG](#).*

## § 4.1 Overview

*This section is non-normative.*

To determine a canonical labeling, [RDPC-1.0](#) considers the information connected to each blank node. Nodes with unique first degree information can immediately be issued a canonical identifier via the [Issue Identifier algorithm](#). When a node has non-unique first degree information, it is necessary to determine all information that is transitively connected to it throughout the entire dataset. [4.6 Hash First Degree Quads](#) defines a node's first degree information via its first degree hash.

Hashes are computed from the information of each blank node. These hashes encode the [mentions](#) incident to each blank node. The [hash](#) of a string  $s$ , is the lower-case, hexadecimal representation of the result of passing  $s$  through a cryptographic hash function. By default, [RDPC-1.0](#) uses the SHA-256 hash algorithm [[FIPS-180-4](#)].

### NOTE

The "degree" terminology is used within this specification as colloquial way of describing the [eccentricity](#) or [radius](#) of any two nodes within a dataset. This concept is also related to "degrees of separation", as in, for example, "[six degrees of separation](#)". Nodes with unique first degree information can be considered nodes with a radius of one.

## § 4.2 Canonicalization State

When performing the steps required by the canonicalization algorithm, it is helpful to track state in a data structure called the ***canonicalization state***. The information contained in the [canonicalization state](#) is described below.

### ***blank node to quads map***

A [map](#) that relates a [blank node identifier](#) to the [quads](#) in which they appear in the [input dataset](#).

### ***hash to blank nodes map***

A [map](#) that relates a [hash](#) to a [list](#) of [blank node identifiers](#).

### ***canonical issuer***

An [identifier issuer](#), initialized with the prefix **c14n** (short for canonicalization), for issuing canonical [blank node identifiers](#).

#### NOTE

Mapping all [blank nodes](#) to use this identifier spec means that an [RDF dataset](#) composed of two different [RDF graphs](#) will issue different identifiers than that for the graphs taken independently. This may happen anyway, due to [automorphisms](#), or overlapping statements, but an identifier based on the resulting [hash](#) along with an issue sequence number specific to that [hash](#) would stand a better chance of surviving such minor changes, and allow the resulting information to be useful for [RDF Diff](#).

## § 4.3 Blank Node Identifier Issuer State

The canonicalization algorithm issues identifiers to [blank nodes](#). The [Issue Identifier algorithm](#) uses an [identifier issuer](#) to accomplish this task. The information an [identifier issuer](#) needs to keep track of is described below.

### ***identifier prefix***

The identifier prefix is a string that is used at the beginning of an [blank node identifier](#). It should be initialized to a string that is specified by the canonicalization algorithm. When generating a new [blank node identifier](#), the prefix is concatenated with a [identifier counter](#). For example, **c14n** is a proper initial value for the [identifier prefix](#) that would produce [blank node identifiers](#) like **c14n1**.

***identifier counter***

A counter that is appended to the [identifier prefix](#) to create an [blank node identifier](#). It is initialized to **0**.

***issued identifiers map***

An [ordered map](#) that relates [blank node identifiers](#) to issued identifiers, to prevent issuance of more than one new identifier per existing identifier, and to allow [blank nodes](#) to be assigned identifiers some time after issuance.

## § 4.4 Canonicalization Algorithm

The canonicalization algorithm converts an [input dataset](#) into a [canonicalized dataset](#) or raises an error if the input dataset is determined to be overly complex. This algorithm will assign deterministic identifiers to any [blank nodes](#) in the [input dataset](#).

### § 4.4.1 Overview

*This section is non-normative.*

[RDPC-1.0](#) canonically labels an [RDF dataset](#) by assigning each [blank node](#) a canonical identifier. In [RDPC-1.0](#), an RDF dataset  $D$  is represented as a set of [quads](#) of the form  $\langle s, p, o, g \rangle$  where the graph component  $g$  is empty if and only if the [triple](#)  $\langle s, p, o \rangle$  is in the [default graph](#). It is expected that, for two RDF datasets, [RDPC-1.0](#) returns the same canonically labeled list of quads if and only if the two datasets are isomorphic (i.e., the same modulo blank node identifiers).

[RDPC-1.0](#) consists of several sub-algorithms. These sub-algorithms are introduced in the following sub-sections. First, we give a high level summary of [RDPC-1.0](#).

**1) Initialization.** Initialize the state needed for the rest of the algorithm using [4.2 Canonicalization State](#). Also initialize the [canonicalized dataset](#) using the [input dataset](#) (which remains immutable) the [input blank node identifier map](#) (retaining blank node identifiers from the input if possible, otherwise assigning them arbitrarily); the [issued identifiers map](#) from the [canonical issuer](#) is added upon completion of the algorithm.

**2) Compute first degree hashes.** Compute the first degree hash for each blank node in the dataset using [4.6 Hash First Degree Quads](#).

**3) Canonically label unique nodes.** Assign canonical identifiers via [4.5 Issue Identifier Algorithm](#), in [Unicode code point order](#), to each blank node whose first degree hash is unique.

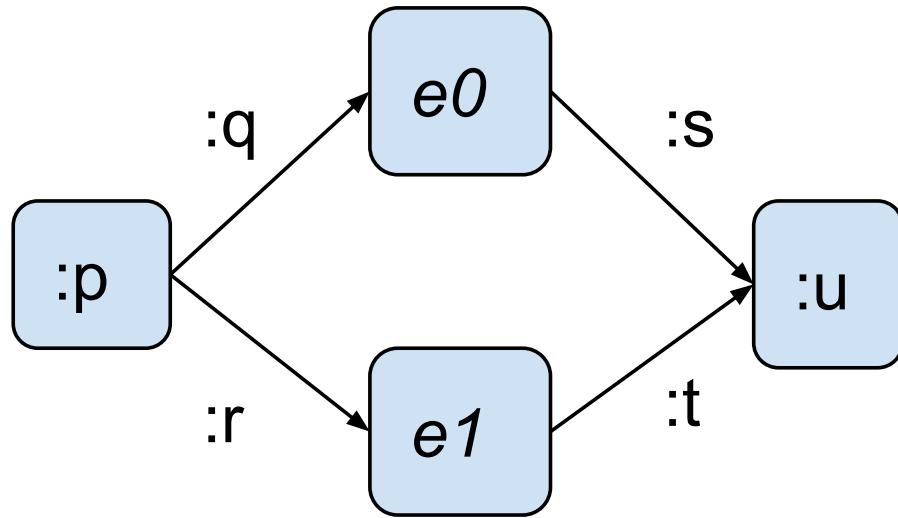
- 4) **Compute N-degree hashes for non-unique nodes.** For each repeated first degree hash (proceeding in [Unicode code point order](#)), compute the N-degree hash via [4.8 Hash N-Degree Quads](#) of every unlabeled blank node that corresponds to the given repeated hash.
- 5) **Canonically label remaining nodes.** In [Unicode code point order](#) of the N-degree hashes, issue canonical identifiers to each corresponding blank node using [4.5 Issue Identifier Algorithm](#). If more than one node produces the same N-degree hash, the order in which these nodes receive a canonical identifier does not matter.
- 6) **Finish.** Return the [serialized canonical form](#) of the [canonicalized dataset](#). Alternatively, return the [canonicalized dataset](#) containing the [input blank node identifier map](#) and [issued identifiers map](#).

#### § 4.4.2 Examples

*This section is non-normative.*

## EXAMPLE 2: Unique hashes

This example illustrates the [Canonicalization Algorithm](#) where all blank nodes have unique first-degree hashes.



*Figure 2 An illustration of a graph resulting in unique hashes.*

*Image available in [SVG](#).*

```

:p :q _:e0 .
:p :r _:e1 .
_:e0 :s :u .
_:e1 :t :u .
  
```

Step 2 is called twice, with each blank node (**e0** and **e1**) in the [input dataset](#) to populate [blank node to quads map](#):

<i>blank node</i>	<i>Q</i>
<b>e0</b>	<b>:p :q e0 .</b> <b>e0 :s :u .</b>
<b>e1</b>	<b>:p :r e1 .</b> <b>e1 :t :u .</b>

*Table 1* Blank node to quads map for unique hashes

Step 3 generates the first-degree hash for each blank node, which is explored further in [Example 4](#):

<i>hash</i>	<i>blank node(s)</i>
21d1dd5ba21f3dee9d76c00c260fa6f5d5d65315099e553026f4828d0dc77a	e0
6fa0b9bdb376852b5743ff39ca4cbf7ea14d34966b2828478fbf222e7c764473	e1

*Table 2* Hash to blank nodes map for unique hashes

Step 4 creates canonical identifiers for each blank node which has a unique hash:

<i>blank node</i>	<i>canonical identifier</i>
e0	c14n0
e1	c14n1

*Table 3* Canonical identifiers for blank nodes with unique hashes

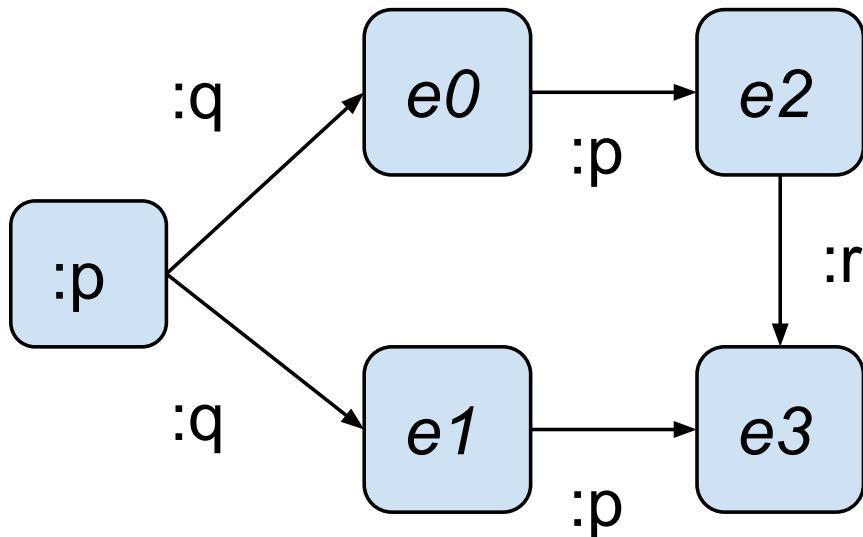
Step 5 has no effect, as there are no remaining blank nodes without canonical identifiers.

Step 6 generates the canonicalized dataset by mapping blank node identifiers in the input dataset with canonical identifiers:

```
:p :q _:c14n0 .
:p :r _:c14n1 .
_:c14n0 :s :u .
_:c14n1 :t :u .
```

### EXAMPLE 3: Shared hashes

This example illustrates the [Canonicalization Algorithm](#) where hashing the statements [mentioning](#) those blank nodes have overlapping results.



[Figure 3](#) An illustration of a graph resulting in shared hashes.

Image available in [SVG](#).

```

:p :q _:e0 .
:p :q _:e1 .
_:e0 :p _:e2 .
_:e1 :p _:e3 .
_:e2 :r _:e3 .
  
```

[Step 2](#) is called four times, with each blank node ([e0](#), [e1](#), [e2](#), and [e3](#)) to populate [blank node to quads map](#):

blank node	Q
e0	<p><a href="#">:p :q e0 .</a></p> <p><a href="#">e0 :p e2 .</a></p>
e1	<p><a href="#">:p :q e1 .</a></p> <p><a href="#">e1 :p e3 .</a></p>
e2	<p><a href="#">e0 :p e2 .</a></p> <p><a href="#">e2 :r e3 .</a></p>

<i>blank node</i>	<i>Q</i>
e3	e1 :p e3 . e2 :r e3 .

*Table 4* Blank node to quads map for shared hashes

[Step 3](#) generates the first-degree hash for each blank node, which is explored further in [Example 5](#) (note that the hashes for *e0* and *e1* are shared):

<i>hash</i>	<i>blank node(s)</i>
3b26142829b8887d011d779079a243bd61ab53c3990d550320a17b59ade6ba36	<i>e0, e1</i>
15973d39de079913dac841ac4fa8c4781c0febfb5e83e5c6e250869587f8659	<i>e2</i>
7e790a99273eed1dc57e43205d37ce232252c85b26ca4a6ff74ff3b5aea7bccd	<i>e3</i>

*Table 5* Hash to blank nodes map shared hashes

[Step 4](#) creates canonical identifiers for each blank node which has a unique hash:

<i>blank node</i>	<i>canonical identifier</i>
e2	c14n0
e3	c14n1

*Table 6* Canonical identifiers for blank nodes with shared hashes

[Step 5](#) is run on *e0* and *e1*, separately, which share the same hash, and use separate instances of a *temporary issuer*, (as explored in [Example 7](#)) to create the *hash path list* composed of the hash result and temporary identifier mappings from [Hash N-Degree Quads algorithm](#) for each of these blank nodes:

<i>temporary issuer mappings</i>	<i>original identifier</i>	<i>temporary identifier</i>
	<i>e0</i>	<i>b0</i>
<i>hash</i>	fbc300de5afafdf97a4b9ee1e72b57754dcdbc7ebb724789ac6a94a5b82a48d30	

*Table 7* Result for *e0*

<i>temporary issuer mappings</i>	<i>original identifier</i>	<i>temporary identifier</i>
	<i>e1</i>	<i>b0</i>
<i>hash</i>	2c0b377baf86f6c18fed4b0df6741290066e73c932861749b172d1e5560f5045	

*Table 8* Result for *e1*

[Step 5.3](#) creates the canonical identifiers for the temporary identifiers established in the previous step running in order of the *hash* component from each result. This updates the [canonical issuer](#) with the following mappings:

<i>blank node</i>	<i>canonical identifier</i>
<i>e0</i>	<i>c14n3</i>
<i>e1</i>	<i>c14n2</i>
<i>e2</i>	<i>c14n0</i>
<i>e3</i>	<i>c14n1</i>

[Table 9](#) Blank node to canonical identifiers

[Step 6](#) updates the [canonicalized dataset](#) with the [canonical issuer](#), containing an [issued identifiers map](#) mapping blank node identifiers from the input dataset to their canonical identifiers:

```
:p :q _:c14n2 .
:p :q _:c14n3 .
_:c14n0 :r _:c14n1 .
_:c14n2 :p _:c14n1 .
_:c14n3 :p _:c14n0 .
```

#### § 4.4.3 Algorithm

The following algorithm will run with a minimal number of iterations in each step for typical [input datasets](#). In some extreme cases, the algorithm can behave poorly, particularly in [Step 5](#).

Implementations *MUST* defend against potential denial-of-service attacks by raising suitable exceptions and terminating early. See [7.1 Dataset Poisoning](#) for further information.

#### NOTE

Implementations can consider placing limits on the number of calls to [4.8 Hash N-Degree Quads](#) based on the number of blank nodes in the [hash to blank nodes map](#). For most typical datasets, more than a couple of iterations on [4.8 Hash N-Degree Quads](#) per blank node would be unusual.

- 1) Create the [canonicalization state](#). If the [input dataset](#) is an N-Quads document, parse that document into a dataset in the [canonicalized dataset](#), retaining any blank node identifiers used

within that document in the [input blank node identifier map](#); otherwise arbitrary identifiers are assigned for each blank node.

► **Explanation**

**2)** For every [quad](#)  $Q$  in [input dataset](#):

**2.1)** For each [blank node](#) that is a component of  $Q$ , add a reference to  $Q$  from the [map entry](#) for the [blank node identifier identifier](#) in the [blank node to quads map](#), creating a new entry if necessary, using the identifier for the blank node found in the [input blank node identifier map](#).

► **Explanation**

► **Logging**

**3)** For each [key](#)  $n$  in the [blank node to quads map](#):

► **Explanation**

**3.1)** Create a [hash](#),  $h_f(n)$ , for  $n$  according to the [Hash First Degree Quads algorithm](#).

**3.2)** Append  $n$  to the value associated to  $h_f(n)$  in [hash to blank nodes map](#), creating a new entry if necessary.

► **Logging**

**4)** For each [hash to identifier list map entry](#) in [hash to blank nodes map](#), [code point ordered](#) by [hash](#):

► **Explanation**

**4.1)** If [identifier list](#) has more than one entry, continue to the next mapping.

**4.2)** Use the [Issue Identifier algorithm](#), passing [canonical issuer](#) and the single [blank node identifier](#), [identifier](#) in [identifier list](#) to issue a canonical replacement identifier for [identifier](#).

**4.3)** Remove the [map entry](#) for [hash](#) from the [hash to blank nodes map](#).

► **Logging**

**5)** For each [hash to identifier list map entry](#) in [hash to blank nodes map](#), [code point ordered](#) by [hash](#):

► **Explanation**

► **Logging**

**5.1)** Create [hash path list](#) where each item will be a result of running the [Hash N-Degree Quads algorithm](#).

► **Explanation**

**5.2)** For each [blank node identifier](#)  $n$  in *identifier list*:

**5.2.1)** If a canonical identifier has already been issued for  $n$ , continue to the next [blank node identifier](#).

**5.2.2)** Create *temporary issuer*, an [identifier issuer](#) initialized with the prefix **b**.

**5.2.3)** Use the [Issue Identifier algorithm](#), passing *temporary issuer* and  $n$ , to issue a new temporary [blank node identifier](#)  $b_n$  to  $n$ .

**5.2.4)** Run the [Hash N-Degree Quads algorithm](#), passing the [canonicalization state](#),  $n$  for *identifier*, and *temporary issuer*, appending the result to the *hash path list*.

► [Logging](#)

**5.3)** For each *result* in the *hash path list*, [code point ordered](#) by the *hash* in *result*:

► [Explanation](#)

**5.3.1)** For each [blank node identifier](#), *existing identifier*, that was issued a temporary identifier by *identifier issuer* in *result*, issue a canonical identifier, in the same order, using the [Issue Identifier algorithm](#), passing [canonical issuer](#) and *existing identifier*.

► [Explanation](#)

► [Logging](#)

**6)** Add the [issued identifiers map](#) from the [canonical issuer](#) to the [canonicalized dataset](#).

► [Explanation](#)

► [Logging](#)

**7)** Return the [serialized canonical form](#) of the [canonicalized dataset](#). Upon request, alternatively (or additionally) return the [canonicalized dataset](#) itself, which includes the [input blank node identifier map](#), and [issued identifiers map](#) from the [canonical issuer](#).

#### NOTE

Technically speaking, one implementation might return a [canonicalized dataset](#) that maps particular blank nodes to different identifiers than another implementation, however, this only occurs when there are isomorphisms in the dataset such that a canonically serialized expression of the dataset would appear the same from either implementation.

► [Explanation](#)

## § 4.5 Issue Identifier Algorithm

This algorithm issues a new [blank node identifier](#) for a given existing [blank node identifier](#). It also updates state information that tracks the order in which new [blank node identifiers](#) were issued. The order of issuance is important for canonically labeling blank nodes that are isomorphic to others in the dataset.

### § 4.5.1 Overview

The algorithm maintains an [issued identifiers map](#) to relate an existing [blank node identifier](#) from the [input dataset](#) to a new [blank node identifier](#) using a given [identifier prefix](#) ([c14n](#)) with new identifiers issued by appending an incrementing number. For example, when called for a blank node identifier such as [e3](#), it might result in a issued identifier of [c14n1](#).

### § 4.5.2 Algorithm

The algorithm takes an [identifier issuer](#)  $I$  and an *existing identifier* as inputs. The output is a new *issued identifier*. The steps of the algorithm are:

- 1) If there is a [map entry](#) for *existing identifier* in [issued identifiers map](#) of  $I$ , return it.
- 2) Generate *issued identifier* by concatenating [identifier prefix](#) with the string value of [identifier counter](#).
- 3) Add an [entry](#) mapping *existing identifier* to *issued identifier* to the [issued identifiers map](#) of  $I$ .
- 4) Increment [identifier counter](#).
- 5) Return *issued identifier*.

## § 4.6 Hash First Degree Quads

This algorithm calculates a [hash](#) for a given [blank node](#) across the [quads](#) in a [dataset](#) in which that blank node is a component. If the hash uniquely identifies that blank node, no further examination is necessary. Otherwise, a hash will be created for the blank node using the algorithm in [4.8 Hash N-Degree Quads](#) invoked via [4.4 Canonicalization Algorithm](#).

### § 4.6.1 Overview

*This section is non-normative.*

To determine whether the first degree information of a node  $n$  is unique, a [hash](#) is assigned to its [mention set](#),  $Q_n$ . The first degree hash of a blank node  $n$ , denoted  $h_f(n)$ , is the hash that results from [4.6 Hash First Degree Quads](#) when passing  $n$ . Nodes with unique first degree hashes have unique first degree information.

For consistency, [blank node identifiers](#) used in  $Q_n$  are replaced with placeholders in a [canonical n-quads](#) serialization of that quad. Every blank node component is replaced with either [a](#) or [z](#), depending on if that component is  $n$  or not.

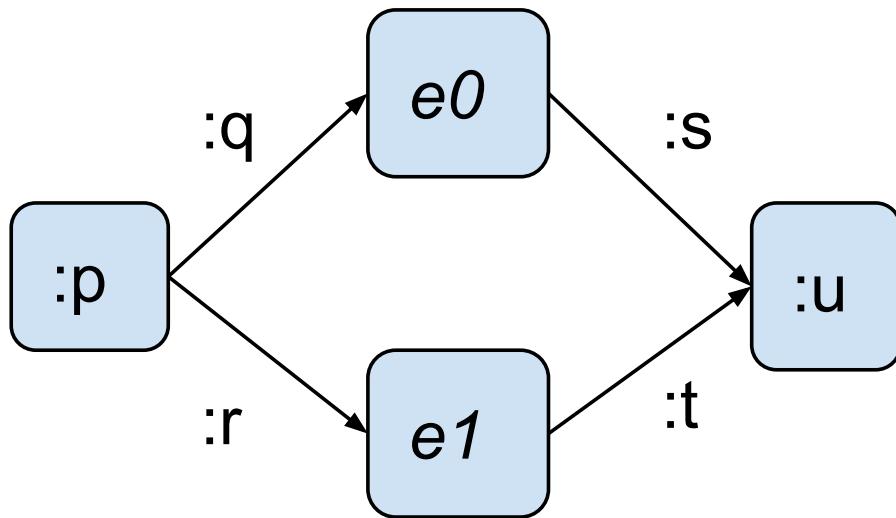
The resulting serialized quads are then [code point ordered](#), concatenated, and hashed. This hash is the first degree hash of  $n$ ,  $h_f(n)$ .

### § 4.6.2 Examples

*This section is non-normative.*

#### EXAMPLE 4: Unique hashes

This example illustrates hashing quads containing blank nodes where hashing the statements mentioning those blank nodes generates unique results.



*Figure 4* An illustration of a graph resulting in unique hashes.

*Image available in [SVG](#).*

```

:p :q _:e0 .
:p :r _:e1 .
_:e0 :s :u .
_:e1 :t :u .
  
```

The algorithm will be called twice, with each blank node (**e0** and **e1**). Running the algorithm with the reference node **e0** results in the following quads, after replacing blank nodes:

```

:p :q _:a .
_:a :s :u .
  
```

These are then serialized to canonical n-quads form: '`<http://example.com/#p>`  
`<http://example.com/#q> _:a .\n_:a <http://example.com/#s>`  
`<http://example.com/#u> .\n'`, concatenated and hashed using the hash algorithm (SHA-256) resulting in `21d1dd5ba21f3dee9d76c0c00c260fa6f5d5d65315099e553026f4828d0dc77a`.

The algorithm is run a second time with the reference node **e1** resulting in the following quads:

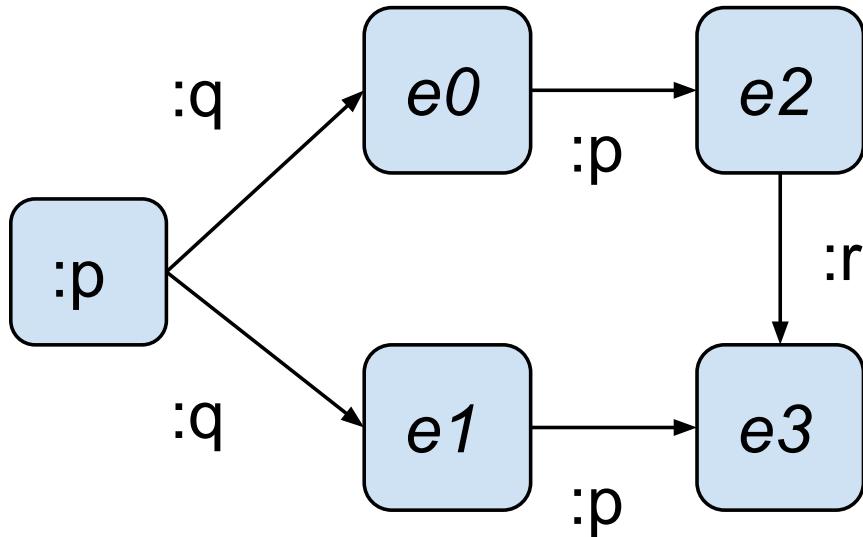
```
:p :r _:a .  
_:a :t :u .
```

These are then serialized to canonical n-quads form: '<http://example.com/#p> <http://example.com/#r> \_:a .\n\_:a <http://example.com/#t> <http://example.com/#u> .\n', concatenated and hashed as before resulting in **6fa0b9bdb376852b5743ff39ca4cbf7ea14d34966b2828478fbf222e7c764473**.

Thus the generated hashes each reference just a single blank node, allowing the canonicalization algorithm to use only the Hash First Degree Quads algorithm.

### EXAMPLE 5: Shared hashes

This example illustrates hashing quads containing blank nodes where hashing the statements [mentioning](#) those blank nodes have overlapping results.



[Figure 5](#) An illustration of a graph resulting in shared hashes.

Image available in [SVG](#).

```

:p :q _:e0 .
:p :q _:e1 .
_:e0 :p _:e2 .
_:e1 :p _:e3 .
_:e2 :r _:e3 .
  
```

The algorithm will be called four times, with each blank node (**e0**, **e1**, **e2**, and **e3**). Running the algorithm with the reference node **e0** results in the following quads, after replacing blank nodes:

```

:p :q _:a .
_:a :p _:z .
  
```

Which hashes to: **3b26142829b8887d011d779079a243bd61ab53c3990d550320a17b59ade6ba36**.

Note that using reference node **e1** results in the same quads, and thus results in the same hash.

Using the reference node **e2** results in the following quads:

```
_:z :p _:a .
_:a :r _:z .
```

Which hashes to: **15973d39de079913dac841ac4fa8c4781c0febfb5e83e5c6e250869587f8659**.

Lastly, using the reference node **e3** results in the following quads:

```
_:z :p _:a .
_:z :r _:a .
```

Which hashes to: **7e790a99273eed1dc57e43205d37ce232252c85b26ca4a6ff74ff3b5aea7bccd**.

The hashes for **e2** and **e3** are unique, but **e0**, and **e1** share a common hash, which will require the use of the [Hash N-Degree Quads Algorithm](#), as it is necessary to consider quads further removed from the direct [mentions](#) to determine a unique hash.

### § 4.6.3 Algorithm

This algorithm takes the [canonicalization state](#) and a **reference blank node identifier** as inputs.

- 1) Initialize **nquads** to an empty [list](#). It will be used to store quads in [canonical n-quads form](#).
- 2) Get the [list](#) of [quads](#) *quads* from the [map entry](#) for [reference blank node identifier](#) in the [blank node to quads map](#).
- 3) For each [quad](#) *quad* in *quads*:
  - 3.1) Serialize the [quad](#) in [canonical n-quads form](#) with the following special rule:
    - 3.1.1) If any component in *quad* is an [blank node](#), then serialize it using a special identifier as follows:
      - 3.1.1.1) If the [blank node](#)'s existing [blank node identifier](#) matches the [reference blank node identifier](#) then use the [blank node identifier](#) **a**, otherwise, use the [blank node identifier](#) **z**.
  - 4) Sort [nquads](#) in [Unicode code point order](#).
  - 5) Return the [hash](#) that results from passing the sorted and concatenated [nquads](#) through the [hash algorithm](#).

► Logging

## § 4.7 Hash Related Blank Node

This algorithm generates a [hash](#) for some [blank node](#) component of a [quad](#), considering its position within that quad. This is used as part of the [Hash N-Degree Quads algorithm](#) to characterize the blank nodes related to some particular blank node within their [mention sets](#).

### § 4.7.1 Overview

*This section is non-normative.*

An *identifier* for a [blank node, related](#), which is a component of some [quad](#), is used along with information describing its position within the quad to create a representation of that blank node within the quad, which is then hashed. The identifier is first found from the following:

- 1) An issued identifier from [canonical issuer](#) in the [canonicalization state](#),
- 2) An issued identifier from *issuer*, or
- 3) the result of the [Hash N-Degree Quads algorithm](#) for *related*.

### § 4.7.2 Examples

*This section is non-normative.*

### EXAMPLE 6: With assigned canonical identifier

This example illustrates generating the related hash for a blank node component of a quad, where the component is in the object position and has already had an assigned canonical identifier from the canonical issuer.

Given the predicate `:p`, related blank node e2 (already assigned the canonical blank node identifier c14n0), an input is created as `o<http://example.com/#p>:_c14n0`, which hashes to `29cf7e22790bc2ed395b81b3933e5329fc7b25390486085cac31ce7252ca60fa`.

If such an identifier had not yet been established, it would instead use the results of the Hash First Degree Quads algorithm, without the leading `_`:

### § 4.7.3 Algorithm

This algorithm creates a hash to identify how one blank node is related to another. It takes the canonicalization state, a related blank node identifier, a quad, an identifier issuer, issuer, and a string position as inputs.

- 1) Initialize a string input to the value of position.
- 2) If position is not g, append <, the value of the predicate in quad, and > to input.
- 3) If there is a canonical identifier for related, or an identifier issued by issuer, append the string \_, followed by that identifier (using the canonical identifier if present, otherwise the one issued by issuer) to input.

► **Explanation**

- 4) Otherwise, append the result of the Hash First Degree Quads algorithm, passing related to input.

► **Explanation**

- 5) Return the hash that results from passing input through the hash algorithm.

► **Explanation**

► **Logging**

## § 4.8 Hash N-Degree Quads

This algorithm calculates a [hash](#) for a given [blank node](#) across the [quads](#) in a [dataset](#) in which that blank node is a component for which the hash does not uniquely identify that blank node. This is done by expanding the search from quads directly referencing that blank node (the [mention set](#)), to those quads which contain nodes which are also components of quads in the mention set, called the [gossip path](#). This process proceeds in every greater degrees of indirection until a unique hash is obtained.

### § 4.8.1 Overview

*This section is non-normative.*

Usually, when trying to determine if two nodes in a graph are equivalent, you simply compare their identifiers. However, what if the nodes don't have identifiers? Then you must determine if the two nodes have equivalent connections to equivalent nodes all throughout the whole graph. This is called the [graph isomorphism](#) problem. This algorithm approaches this problem by considering how one might draw a graph on paper. You can test to see if two nodes are equivalent by drawing the graph twice. The first time you draw the graph the first node is drawn in the center of the page. If you can draw the graph a second time such that it looks just like the first, except the second node is in the center of the page, then the nodes are equivalent. This algorithm essentially defines a deterministic way to draw a graph where, if you begin with a particular node, the graph will always be drawn the same way. If two graphs are drawn the same way with two different nodes, then the nodes are equivalent. A [hash](#) is used to indicate a particular way that the graph has been drawn and can be used to compare nodes.

When two blank nodes have the same first degree hash, extra steps must be taken to detect global, or  $N$ -degree, distinctions. All information that is in any way connected to the blank node  $n$  through other blank nodes, even transitively, must be considered.

To consider all transitive information, the algorithm traverses and encodes all possible paths of incident [mentions](#) emanating from  $n$ , called [gossip paths](#), that reach every unlabeled blank node connected to  $n$ . Each unlabeled blank node is assigned a temporary identifier in the order in which it is reached in the gossip path being explored. The mentions that are traversed to reach connected blank nodes are encoded in these paths via related hashes. This provides a deterministic way to order all

paths coming from  $n$  that reach all blank nodes connected to  $n$  without relying on input blank node identifiers.

This algorithm works in concert with the main canonicalization algorithm to produce a unique, deterministic identifier for a particular blank node. This [hash](#) incorporates all of the information that is connected to the blank node as well as how it is connected. It does this by creating deterministic paths that emanate out from the blank node through any other adjacent blank nodes.

Ultimately, the algorithm selects the shortest [gossip path](#) (based on its encoding as a string), distributing canonical identifiers to the unlabeled blank nodes in the order in which they appear in this path. The hash of this encoded shortest path, called the N-degree hash of  $n$ , distinguishes  $n$  from other blank nodes in the dataset.

For clarity, we consider a [gossip path](#) encoded via the string  $s$  to be shortest provided that:

- 1)** The length of  $s$  is less than or equal to the length of any other gossip path string  $s'$ .
- 2)** If  $s$  and  $s'$  have the same length (as strings), then  $s$  is [code point ordered](#) less than or equal to  $s'$ .

For example,  $abc$  is shorter than  $bbc$ , whereas  $abcd$  is longer than  $bcd$ .

The following provides a high level outline for how the N-degree hash of  $n$  is computed along the shortest [gossip path](#). Note that the full algorithm considers all gossip paths, ultimately returning the [hash](#) of the shortest encoded path.

**1) Compute related hashes.** Compute the related hash  $H_n$  set for  $n$ , i.e., all first degree [mentions](#) between  $n$  and another blank node. Note that this includes both unlabeled blank nodes and those already issued a canonical identifier (labeled blank nodes).

**2) Explore mentions.** Given the related hash  $x$  in  $H_n$ , record  $x$  in the data to hash  $D_n$ . Determine whether each blank node reachable via the [mention](#) with related hash  $x$  has already received an identifier.

**2.1) Record the identifiers of labeled nodes.** If a blank node already has an identifier, record its identifier in  $D_n$  once for every [mention](#) with related hash  $x$ . Skip to the next related hash in  $H_n$  and repeat [step 2](#).

**2.2) Distribute and record temporary identifiers to unlabeled nodes.** For each unlabeled blank node, assign it a temporary identifier according to the order in which it is reached in the gossip path, recording its given identifier in  $D_n$  (including repetitions). Add each unlabeled node to the recursion list  $R_n(x)$  in this same order (omitting repetitions).

**2.3) Recurse on newly labeled nodes.** For each  $n_i$  in  $R_n(x)$

**2.3.1)** Record its identifier in  $D_n$

**2.3.2)** Append  $\langle r(i) \rangle$  to  $D_n$  where  $r(i)$  is the data to hash that results from returning to [step 1](#), replacing  $n$  with  $n_i$ .

**3) Compute the  $N$ -degree hash of  $n$ .** Hash  $D_n$  to return the  $N$ -degree hash of  $n$ , namely  $h_N(n)$ .

Return the updated issuer  $I_n$  that has now distributed temporary identifiers to all unlabeled blank nodes connected to  $n$ .

As described above in [step 2.3](#),  $HN$  recurses on each unlabeled blank node when it is first reached along the [gossip path](#) being explored. This recursion can be visualized as moving along the path from  $n$  to the blank node  $n_i$  that is receiving a temporary identifier. If, when recursing on  $n_i$ , another unlabeled blank node  $n_j$  is discovered, the algorithm again recurses. Such a recursion traces out the [gossip path](#) from  $n$  to  $n_j$  via  $n_i$ .

The recursive hash  $r(i)$  is the [hash](#) returned from the completed recursion on the node  $n_i$  when computing  $h_N(n)$ . Just as  $h_N(n)$  is the hash of  $D_n$ , we denote the data to hash in the recursion on  $n_i$  as  $D_i$ . So,  $r(i) = h(D_i)$ . For each related hash  $x \in H_n$ ,  $R_n(x)$  is called the *recursion list* on which the algorithm recurses.

## § 4.8.2 Examples

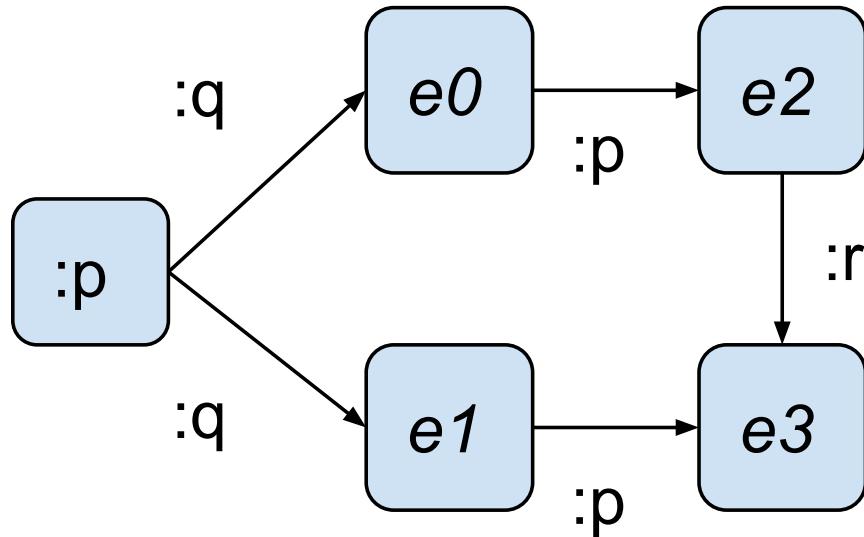
*This section is non-normative.*

### EXAMPLE 7: Shared hashes

This example revisits [Example 3](#) to illustrate the operation of the [Hash N-Degree Quads Algorithm](#) on the blank nodes (`e0` and `e1`) which resulted in a shared result from the [Hash First Degree Quads Algorithm](#). The other blank nodes (`e2` and `e3`) had unique hashes when processed by the [Hash N-Degree Quads Algorithm](#), so that canonical identifiers have already been issued, as recorded in the [identifier issuer](#) instance.

<i>blank node</i>	<i>canonical identifier</i>
<code>e2</code>	<code>c14n0</code>
<code>e3</code>	<code>c14n1</code>

*Table 10* Established canonical identifiers from [canonical issuer](#)



*Figure 6* An illustration of a graph resulting in shared hashes.

*Image available in [SVG](#).*

```

:p :q _:e0 .
:p :q _:e1 .
_:e0 :p _:e2 .
_:e1 :p _:e3 .
_:e2 :r _:e3 .
  
```

The algorithm will be called twice, with each blank node which did not result in a unique hash from the [Hash First Degree Quads Algorithm](#) (`e0` and `e1`). The [map entry](#) for `e0` in the [blank node to quads map](#) results in the following quads:

```
:p :q _:e0 .
_:e0 :p _:e2 .
```

When called, the *temporary issuer* has the following mappings:

<b>blank node</b>	<b>temporary identifier</b>
<code>e0</code>	<code>b0</code>

[Table 11](#) Initial mappings in temporary issuer for `e0`

[Step 3](#) iterates over each of these quads to find [blank node](#) components to populate  $H_n$  with the result of running the [Hash Related Blank Node algorithm](#) using the *position* of that *component* within the quad.

<b>related</b>	<b>quad</b>	<b>position</b>	<b>hash</b>
<code>e2</code>	<code>e0 :p e2 .</code>	<code>o</code>	<code>29cf7e22790bc2ed395b81b3933e5329fc7b25390486085cac31ce7252ca60fa</code>

[Table 12](#) Hash for `e2` related to `e0`

which results in the following  $H_n$ :

<b>related hash</b>	<b>blank node list</b>
<code>29cf7e22790bc2ed395b81b3933e5329fc7b25390486085cac31ce7252ca60fa</code>	<code>[ e2 ]</code>

[Table 13](#) Blank node list for hash

[Step 5](#) iterates over each *related hash* which is a key in  $H_n$ , which can have one or more related blank nodes, which determines the shortest [gossip path](#) between the two nodes (`e0` and `e2`). In this case, the hash maps to just `e2`, for which a canonical identifier has already been chosen (`c14n0`), so there is a single permutation resulting in one candidate path. The resulting *chosen path* is `_:c14n0`.

The string *data to hash* is composed of the single *related hash* and the *chosen path*:

`29cf7e22790bc2ed395b81b3933e5329fc7b25390486085cac31ce7252ca60fa_:c14n0`, which hashes to `fbc300de5afafdf97a4b9ee1e72b57754dcdb7ebb724789ac6a94a5b82a48d30`, which, along with the *temporary issuer* used to traverse these paths, is the result of the algorithm:

<i>temporary issuer mappings</i>	<i>original identifier</i>	<i>temporary identifier</i>
	<i>e0</i>	<i>b0</i>
<i>hash</i>	<i>fbc300de5afafdf97a4b9ee1e72b57754dcdbc7ebb724789ac6a94a5b82a48d30</i>	

*Table 14 Result for e0*

The algorithm is called again for *e1*.

<i>blank node</i>	<i>temporary identifier</i>
<i>e1</i>	<i>b0</i>

*Table 15 Initial mappings in temporary issuer for e1*

The map entry for *e1* in the blank node to quads map results in the following quads:

```
:p :q _:e1 .
_:e1 :p _:e3 .
```

Step 3 again calculates blank node components using Hash Related Blank Node algorithm.

<i>related</i>	<i>quad</i>	<i>position</i>	<i>hash</i>
<i>e3</i>	<i>e1 :p e3 .</i>	<i>o</i>	<i>b7956ea1d6...ff6b098216</i>

*Table 16 Hash for e3 related to e1*

which results in the following  $H_n$ :

<i>related hash</i>	<i>blank node list</i>
<i>b7956ea1d654d5824496eb439a1f2b79478bd7d02d4a115f4c97cbff6b098216</i>	<i>[ e3 ]</i>

*Table 17 Blank node list for hash*

Step 5 runs in essentially the same manner, mapping just *e3* having the canonical identifier *c14n1*, so there is again a single permutation resulting in one candidate path. The resulting *chosen path* is *\_:c14n1*.

The string *data to hash* is composed of the single *related hash* and the *chosen path*:

*af54b9512b1ef069205e8e41bc5a96e86a108b0389caa5029f2c3fd0bc465246\_:c14n1*, which hashes to *767a5e66f509221f45003a16c12a89d4d9675cfa51ffa80459b63606bdfc2ada*.

The string *data to hash* is composed of the single *related hash* and the *chosen path*:

*b7956ea1d654d5824496eb439a1f2b79478bd7d02d4a115f4c97cbff6b098216\_:c14n1*, which

hashes to `fbc300de5afafdf97a4b9ee1e72b57754dcdb7ebb724789ac6a94a5b82a48d30`, which, along with the *temporary issuer* used to traverse these paths, is the result of the algorithm:

<i>temporary issuer mappings</i>	<i>original identifier</i>	<i>temporary identifier</i>
	<code>e1</code>	<code>b0</code>
<i>hash</i>	<code>2c0b377baf86f6c18fed4b0df6741290066e73c932861749b172d1e5560f5045</code>	

*Table 18 Result for e1*

### § 4.8.3 Algorithm

The inputs to this algorithm are the [canonicalization state](#), the *identifier* for the [blank node](#) to recursively hash quads for, and path identifier *issuer* which is an [identifier issuer](#) that issues temporary [blank node identifiers](#). The output from this algorithm will be a [hash](#) and the [identifier issuer](#) used to help generate it.

#### ► Logging

- 1) Create a new [map](#)  $H_n$  for relating hashes to related [blank nodes](#).
- 2) Get a reference, *quads*, to the list of [quads](#) from the [map entry](#) for *identifier* in the [blank node to quads map](#).

#### ► Explanation

#### ► Logging

- 3) For each *quad* in *quads*:

#### ► Explanation

- 3.1) For each *component* in *quad*, where *component* is the [subject](#), [object](#), or [graph name](#), and it is a [blank node](#) that is not identified by *identifier*:

3.1.1) Set *hash* to the result of the [Hash Related Blank Node algorithm](#), passing the [blank node identifier](#) for *component* as *related*, *quad*, *issuer*, and *position* as either `s`, `o`, or `g` based on whether *component* is a [subject](#), [object](#), [graph name](#), respectively.

3.1.2) Add a mapping of *hash* to the [blank node identifier](#) for *component* to  $H_n$ , adding an entry as necessary.

#### ► Logging

- 4) Create an empty string, *data to hash*.

**5)** For each *related hash* to *blank node list* mapping in  $H_n$ , code point ordered by *related hash*:

► **Explanation**

► **Logging**

**5.1)** Append the *related hash* to the *data to hash*.

**5.2)** Create a string chosen path.

**5.3)** Create an unset *chosen issuer* variable.

**5.4)** For each permutation  $p$  of *blank node list*:

► **Logging**

**5.4.1)** Create a copy of *issuer*, *issuer copy*.

**5.4.2)** Create a string path.

**5.4.3)** Create a *recursion list*, to store blank node identifiers that must be recursively processed by this algorithm.

**5.4.4)** For each *related* in  $p$ :

**5.4.4.1)** If a canonical identifier has been issued for *related* by canonical issuer, append the string \_:, followed by the canonical identifier for *related*, to *path*.

► **Explanation**

**5.4.4.2)** Otherwise:

**5.4.4.2.1)** If *issuer copy* has not issued an identifier for *related*, append *related* to *recursion list*.

► **Explanation**

**5.4.4.2.2)** Use the Issue Identifier algorithm, passing *issuer copy* and the *related*, and append the string \_:, followed by the result, to *path*.

**5.4.4.3)** If *chosen path* is not empty and the length of *path* is greater than or equal to the length of *chosen path* and *path* is greater than *chosen path* when considering code point order, then skip to the next permutation  $p$ .

► **Explanation**

► **Explanation**

► **Logging**

**5.4.5)** For each *related* in *recursion list*:

► **Explanation**

► **Logging**

**5.4.5.1)** Set *result* to the result of recursively executing the [Hash N-Degree Quads algorithm](#), passing the [canonicalization state](#), *related* for *identifier*, and *issuer copy* for *path identifier issuer*.

► [Logging](#)

**5.4.5.2)** Use the [Issue Identifier algorithm](#), passing *issuer copy* and *related*; append the string `_:`, followed by the result, to *path*.

**5.4.5.3)** Append `<`, the [hash](#) in *result*, and `>` to *path*.

**5.4.5.4)** Set *issuer copy* to the [identifier issuer](#) in *result*.

**5.4.5.5)** If *chosen path* is not empty and the length of *path* is greater than or equal to the length of *chosen path* and *path* is greater than *chosen path* when considering [code point order](#), then skip to the next *p*.

► [Explanation](#)

**5.4.6)** If *chosen path* is empty or *path* is less than *chosen path* when considering [code point order](#), set *chosen path* to *path* and *chosen issuer* to *issuer copy*.

**5.5)** Append *chosen path* to *data to hash*.

► [Logging](#)

**5.6)** Replace *issuer*, by reference, with*chosen issuer*.

**6)** Return *issuer* and the [hash](#) that results from passing *data to hash* through the [hash algorithm](#).

► [Logging](#)

## § 5. Serialization

This section describes the process of creating a serialized [N-Quads] representation of a [canonicalized dataset](#).

The *serialized canonical form* of a [canonicalized dataset](#) is an N-Quads document [N-QUADS] created by representing each [quad](#) from the [canonicalized dataset](#) in [canonical n-quads form](#), sorting them into [code point order](#), and concatenating them. (Note that each canonical N-Quads statement ends with a new line, so no additional separators are needed in the concatenation.) The resulting document has a media type of `application/n-quads`, as described in [C. N-Quads Internet Media Type, File Extension and Macintosh File Type](#) of [N-QUADS].

When serializing quads in [canonical n-quads form](#), components which are [blank nodes](#) *MUST* be serialized using the canonical label associated with each [blank node](#) from the [issued identifiers map](#)

component of the [canonicalized dataset](#).

#### EXAMPLE 8: Canonical N-Quads representation of Unique Hashes

This example illustrates the result of serializing the [canonicalized dataset](#) described in [Example 2](#).

```
<http://example.com/p> <http://example.com/q> _:c14n0 .  
<http://example.com/p> <http://example.com/r> _:c14n1 .  
_:c14n0 <http://example.com/s> <http://example.com/u> .  
_:c14n1 <http://example.com/t> <http://example.com/u> .
```

## § 6. Privacy Considerations

*This section is non-normative.*

The nature of the canonicalization algorithm inherently correlates its output, i.e., the canonical labels and the sorted order of quads, with the input dataset. This could pose issues, particularly when dealing with datasets containing personal information. For example, even if certain information is removed from the [canonicalized dataset](#) for some privacy-respecting reason, there remains the possibility that a third party could infer the omitted data by analyzing the canonicalized dataset. If it is necessary to decouple the canonicalization algorithm's input and output, some suitable post-processing methods for the output of the canonicalization should be performed. This specification has been designed to help make additional processing easier, but other specifications that build on top of this one are responsible for providing any specific details. See [Selective Disclosure](#) in [Verifiable Credential Data Integrity 1.0](#) [VC-DATA-INTEGRITY] for more details about such post-processing methods.

## § 7. Security Considerations

*This section is non-normative.*

### § 7.1 Dataset Poisoning

*This section is non-normative.*

The canonicalization algorithm examines every difference in the information connected to blank nodes in order to ensure that each will properly receive its own canonical identifier. This process can be exploited by attackers to construct datasets which are known to take large amounts of computing time to canonicalize, but that do not express useful information or express it using unnecessary complexity. Implementers of the algorithm are expected to add mitigations that will, by default, abort canonicalizing problematic inputs.

Suggested mitigations include, but are not limited to:

- providing a configurable timeout with a default value applicable to an implementation's common use
- providing a configurable limit on the number of iterations of steps performed in the algorithm, particularly recursive steps and permutations of long lists

Additionally, software that uses implementations of the algorithm can employ best-practice schema validation to reject data that does not meet application requirements, thereby preventing useless poison datasets from being processed. However, such mitigations are application specific and not directly applicable to implementers of the canonicalization algorithm itself.

## § 7.2 Insecure Hash Algorithms

*This section is non-normative.*

It is possible that the default [hash algorithm](#) used by [RDFC-1.0](#) might become insecure at some point in the future. To mitigate this, this algorithm and implementations of it can be parameterized to use a different hash function, without the need to make any changes to the canonicalization algorithm itself. However, using a different hash algorithm will generally lead to different results; applications making use of this specification should carefully weigh the advantages and disadvantages of using an alternative hash function.

### NOTE

The possible implications of the default hash algorithm becoming insecure are mitigated by that fact that no internal hash values are revealed, and the canonicalization algorithm is designed to cope with first-degree hash collisions.

## § 8. Use Cases

*This section is non-normative.*

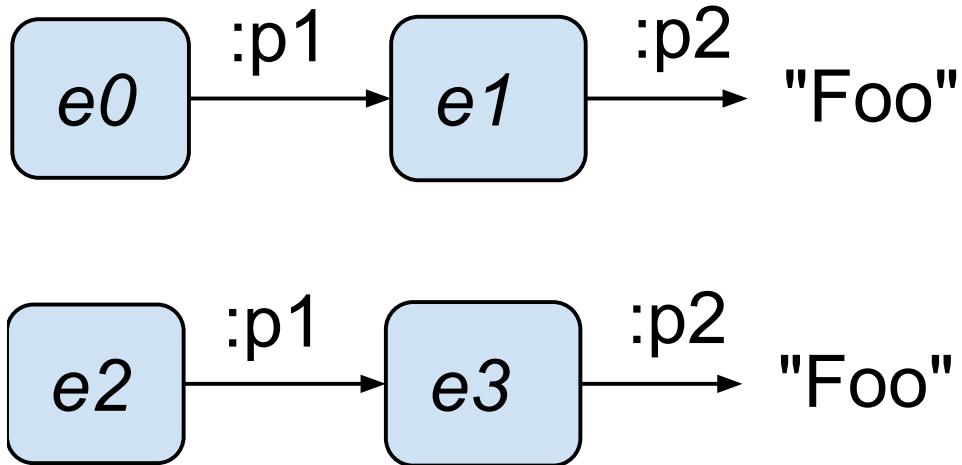
The use cases that have driven the development of the RDF Dataset Canonicalization algorithm are documented in a separate document. It includes further background and explanations for the design decisions taken [[RCH-EXPLAINER](#)].

## § 9. Examples

*This section is non-normative.*

### § 9.1 Duplicate Paths

This example illustrates a more complicated example where the same paths through blank nodes are duplicated in a graph, but use different [blank node identifiers](#).



[Figure 7](#) An illustration of a graph with duplicated paths.

Image available in [SVG](#).

```
_:e0 :p1 _:e1 .  
_:e1 :p2 "Foo" .  
_:e2 :p1 _:e3 .  
_:e3 :p2 "Foo" .
```

The following is a summary of the more detailed execution log found [here](#).

### EXAMPLE 9: Duplicate Paths

Step 2 of the [Canonicalization Algorithm](#) is called four times, with each blank node ( $e0$ ,  $e1$ ,  $e2$ , and  $e3$ ) to populate [blank node to quads map](#):

<i>blank node</i>	<i>Q</i>
$e0$	$_:{e0} :p1 _:e1 .$
$e1$	$_:{e0} :p1 _:e1 .$ $_:{e1} :p2 "Foo" .$
$e2$	$_:{e2} :p1 _:e3 .$
$e3$	$_:{e2} :p1 _:e3 .$ $_:{e3} :p2 "Foo" .$

[Table 19](#) Blank node to quads map for duplicate paths

Step 3 generates the first-degree hash for each blank node.

For  $e0$ , the [Hash First Degree Quads algorithm](#) uses the nquad  $_:{a} :p1 _:z .$  to calculate the hash [24da9a4406b4e66dfffa10ad3d4d6dddc388fbf193bb124e865158ef419893957](#).

For  $e1$ , the [Hash First Degree Quads algorithm](#) uses the nquads  $_:{z} :p1 _:a .$  and  $_:{a} :p2 "Foo" .$  to calculate the hash [a994e40b576809985bc0f389308cd9d552fd7c89d028c163848a6b2d33a8583a](#).

For  $e2$ , the [Hash First Degree Quads algorithm](#) uses the nquad  $_:{a} :p1 _:z .$  to calculate the hash [24da9a4406b4e66dfffa10ad3d4d6dddc388fbf193bb124e865158ef419893957](#).

For  $e3$ , the [Hash First Degree Quads algorithm](#) uses the nquads  $_:{z} :p1 _:a .$  and  $_:{a} :p2 "Foo" .$  to calculate the hash [a994e40b576809985bc0f389308cd9d552fd7c89d028c163848a6b2d33a8583a](#).

<i>hash</i>	<i>blank node(s)</i>
<a href="#">24da9a4406b4e66dfffa10ad3d4d6dddc388fbf193bb124e865158ef419893957</a>	$e0$ and $e2$
<a href="#">a994e40b576809985bc0f389308cd9d552fd7c89d028c163848a6b2d33a8583a</a>	$e1$ and $e3$

[Table 20](#) Hash to blank nodes map shared hashes

Step 4 would create canonical identifiers for each blank node which has a unique hash, but no blank node has a unique hash.

Step 5 is run on  $e_0$ ,  $e_1$ ,  $e_2$ , and  $e_3$ , separately, which share hash values. each use separate instances of a *temporary issuer*, to create the *hash path list* composed of the hash result and temporary identifier mappings from [Hash N-Degree Quads algorithm](#) for each of these blank nodes/

The [Hash N-Degree Quads algorithm](#) is first called for hash  $24da9a4406b4e66dfffa10ad3d4d6dddc388fbf193bb124e865158ef419893957$  and blank node  $e_0$  related to the nquad  $_:{e_0} :p1 _:{e_1} ..$

<b>blank node temporary identifier</b>	
$e_0$	$b_0$

Step 3 calls the [Hash Related Blank Node algorithm](#) for each position in this quad:

<i>related</i>	<i>quads</i>	<i>position</i>	<i>hash</i>
$e_3$	$z :p1 a .$ $a :p2 "Foo" .$	0	$a994e40b57...2d33a8583a$

*Table 21* Hash for  $e_2$  related to  $e_0$

which results in the following  $H_n$ :

<i>related hash</i>	<i>blank node list</i>
$a994e40b576809985bc0f389308cd9d552fd7c89d028c163848a6b2d33a8583a$	[ $e_3$ ]

*Table 22* Blank node list for hash

Step 5 iterates over each *related hash* which is a key in  $H_n$ , which can have one or more related blank nodes.

At Step 5.4.5, the *recursion list* is [  $e_3$  ], so the [Hash N-Degree Quads algorithm](#) is called recursively resulting in the following:

<i>temporary issuer mappings</i>	<i>original identifier</i>	<i>temporary identifier</i>
	$e_2$	$b_0$
$e_3$		$b_1$
<i>hash</i>	$c484f98e6cbf9e21f287433c8b1caa7f1486fd61d84ab220a494bf8184751b8c$	

*Table 23* Result for  $e_3$

Back in [step 5.4.5.4](#), the *path* and *issuer copy* are:

`_ :b1_ :b1<...>` and `{e2: b0, e3: b1}`.

This results in the *chosen path*

`_ :b1_ :b1<...>` and *data to hash*

`3d96946f27fc34a78e8d067135a1cb1b77083aebc4b2c6cbdc536f067242686c_ :b1_ :b1<...>`

<i>temporary issuer mappings</i>	<i>original identifier</i>	<i>temporary identifier</i>
	<code>e2</code>	<code>b0</code>
	<code>e3</code>	<code>b1</code>
<i>hash</i>	<code>39d609fc...409d7</code>	

[Table 24](#) Result for `e2`

The [Hash N-Degree Quads algorithm](#) is next called for the same hash

`24da9a4406b4e66dff...5ef419893957` but this time with blank node `e2` related to the nquad `_ :e2 :p1 _ :e3 ..`

<i>blank node temporary identifier</i>	
<code>e2</code>	<code>b0</code>

[Step 3](#) calls the [Hash Related Blank Node algorithm](#) for each position in this quad:

<i>related</i>	<i>quads</i>	<i>position</i>	<i>hash</i>
<code>e3</code>	<code>z :p1 a .</code> <code>a :p2 "Foo" .</code>	<code>o</code>	<code>3d96946f27....184751b8c</code>

[Table 25](#) Hash for `e2` related to `e0`

which results in the following  $H_n$ :

<i>related hash</i>	<i>blank node list</i>
<code>3d96946f27fc34a78e8d067135a1cb1b77083aebc4b2c6cbdc536f067242686c</code>	<code>[ e3 ]</code>

[Table 26](#) Blank node list for hash

[Step 5](#) iterates over each *related hash* which is a key in  $H_n$ , which can have one or more related blank nodes.

<i>temporary issuer mappings</i>	<i>original identifier</i>	<i>temporary identifier</i>
	<i>e0</i>	<i>b0</i>
<i>hash</i>	<i>fbc300de5afafdf97a4b9ee1e72b57754dcdbc7ebb724789ac6a94a5b82a48d30</i>	

*Table 27 Result for e0*

Step 5.3 back in the [Canonicalization Algorithm](#) creates canonical identifiers for the temporary identifiers just issued:

<i>blank node canonical identifier</i>	
<i>e0</i>	<i>c14n0</i>
<i>e1</i>	<i>c14n1</i>
<i>e2</i>	<i>c14n2</i>
<i>e3</i>	<i>c14n3</i>

Next, back in step 5.1, now using hash

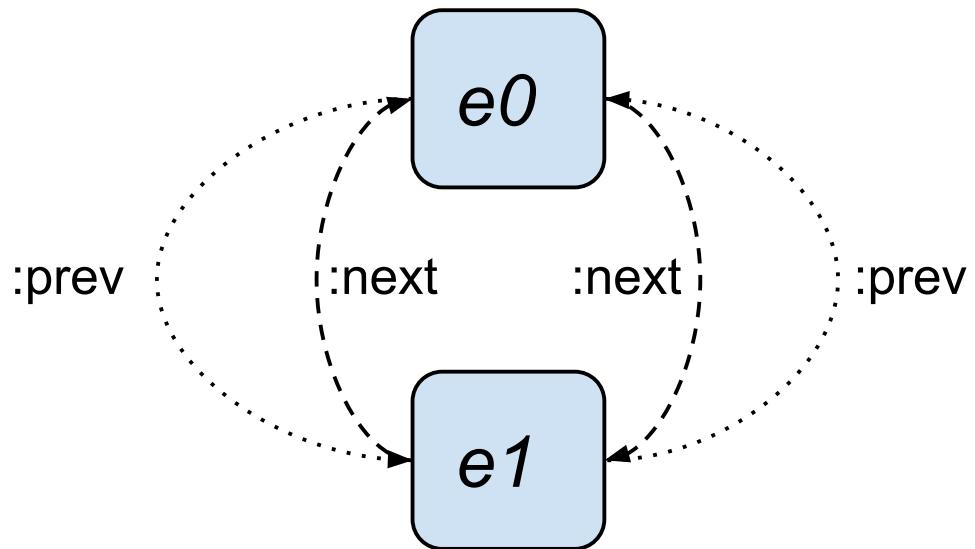
*a994e40b576809985bc0f389308cd9d552fd7c89d028c163848a6b2d33a8583a*, canonical identifiers have already been created for the two blank nodes *e1* and *e3*, so no further processing is necessary.

Step 6 ends with the issued identifiers map containing the following mappings:

<i>blank node canonical identifier</i>	
<i>e0</i>	<i>c14n0</i>
<i>e1</i>	<i>c14n1</i>
<i>e2</i>	<i>c14n2</i>
<i>e3</i>	<i>c14n3</i>

## § 9.2 Double Circle

This example illustrates another complicated example of nodes that are doubly connected in opposite directions.



*Figure 8 An illustration of a graph back and forth links to nodes.*

*Image available in [SVG](#).*

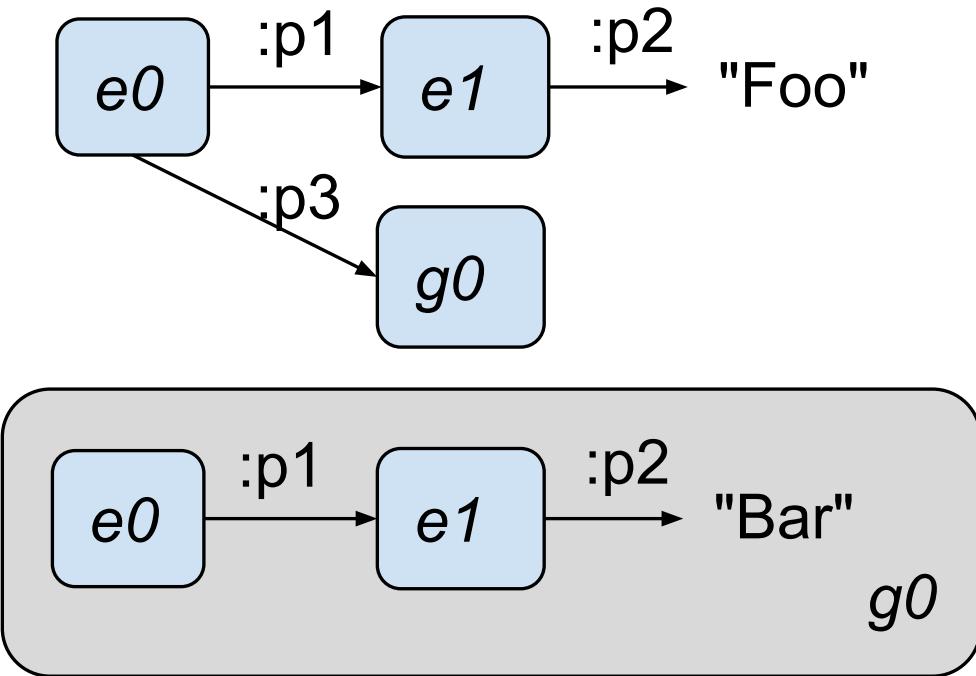
```

_:e0 :next _:e1 .
_:e0 :prev _:e1 .
_:e1 :next _:e0 .
_:e1 :prev _:e0 .
  
```

The example is not explored in detail, but the execution log found [here](#) shows examples of more complicated pathways through the algorithm

### § 9.3 Dataset with Blank Node Named Graph

This example illustrates an example of a [dataset](#), where one graph is named using a [blank node](#), which is also the object of a triple in the [default graph](#).



[Figure 9](#) An illustration of a dataset containing a graph named with a blank node.

Image available in [SVG](#).

```
_:e0 :p1 _:e1 .
_:e1 :p2 "Foo" .
_:e1 :p3 _:g0 .
_:e0 :p1 _:e1 _:g0 .
_:e1 :p2 "Bar" _:g0 .
```

The following is a summary of the more detailed execution log found [here](#).

### EXAMPLE 10: Dataset Blank Node Graph

Step 2 of the [Canonicalization Algorithm](#) is called three times, with each blank node ( $e0$ ,  $e1$ , and  $g0$ ) to populate [blank node to quads map](#):

blank node	$Q$
$e0$	$_e0 :p1 _e1 .$ $_e0 :p1 _e1 _g0 .$
$e1$	$_e0 :p1 _e1 .$ $_e1 :p2 "Foo" .$ $_e1 :p3 _g0 .$ $_e0 :p1 _e1 _g0 .$ $_e1 :p2 "Bar" _g0 .$
$g0$	$_e1 :p3 _g0 .$ $_e0 :p1 _e1 _g0 .$ $_e1 :p2 "Bar" _g0 .$

*Table 28* Blank node to quads map for dataset with blank node named graph

Step 3 generates the first-degree hash for each blank node.

For  $e0$ , the [Hash First Degree Quads algorithm](#) uses the nquads  $_a :p1 _z .$  and  $_a :p1 _z _z .$  to calculate the hash

`39fa92e7a9ea0b32e3e08aeadebc1cbbcd9ba9945a18f2594677598fcfb517345`.

For  $e1$ , the [Hash First Degree Quads algorithm](#) uses the nquads  $_z :p1 _a .$ ,  $_z :p2 "Foo" .$ ,  $_z :p3 _z .$ ,  $_z :p1 _a _z .$ , and  $_z :p2 "Bar" _z .$  to calculate the hash `5274652aa97d22f9f2998d1837ffaf9ad820fb813574adc4a403f105021e48bf`.

For  $g0$ , the [Hash First Degree Quads algorithm](#) uses the nquads  $_z :p3 _a .$ ,  $_z :p1 _z _a .$ , and  $_z :p2 "Bar" _a .$  to calculate the hash `efb0096125247801199082c6f213cda5377d8ddac20a83eaf6c7d88335c12bf4`.

hash	blank node(s)
<code>39fa92e7a9ea0b32e3e08aeadebc1cbbcd9ba9945a18f2594677598fcfb517345</code>	$e0$
<code>5274652aa97d22f9f2998d1837ffaf9ad820fb813574adc4a403f105021e48bf</code>	$e1$
<code>efb0096125247801199082c6f213cda5377d8ddac20a83eaf6c7d88335c12bf4</code>	$g0$

*Table 29* Hash to blank nodes map shared hashes

Step 4 creates canonical identifiers for each blank node which has a unique hash:

<i>blank node</i>	<i>canonical identifier</i>
<i>e0</i>	<i>c14n0</i>
<i>e1</i>	<i>c14n1</i>
<i>g0</i>	<i>c14n2</i>

*Table 30* Canonical identifiers for dataset with blank node named graph

Step 5 has no effect, as there are no remaining blank nodes without canonical identifiers.

Step 6 ends with the canonical issuers containing the following mappings:

<i>blank node canonical identifier</i>	
<i>e0</i>	<i>c14n0</i>
<i>e1</i>	<i>c14n1</i>
<i>g0</i>	<i>c14n2</i>

## § A. A Canonical form of N-Quads

This section defines a canonical form of N-Quads which has a completely specified layout. The grammar for the language remains unchanged.

Canonical N-Quads updates and extends [Canonical N-Triples](#) in [\[N-TRIPLES\]](#) to include [graphLabel](#).

While the N-Quads syntax [\[N-QUADS\]](#) allows choices for the representation and layout of RDF data, the canonical form of N-Quads provides a unique syntactic representation of any quad. Each code point can be represented by only one of [UCHAR](#), [ECHAR](#), or unencoded character, where the relevant production allows for a choice in representation. Each quad is represented entirely on a single line with specified white space.

Canonical N-Quads has the following additional constraints on layout:

- White space *MUST NOT* be used except after [subject](#), [predicate](#), [object](#), and [graphLabel](#), each of which *MUST* be a single space (code point [U+0020](#)).

- Literals with the datatype <http://www.w3.org/2001/XMLSchema#string> *MUST NOT* use the datatype IRI part of the literal, and are represented using only STRING\_LITERAL\_QUOTE.
- HEX *MUST* use only digits ([0-9]) and uppercase letters ([A-F]).
- Within STRING\_LITERAL\_QUOTE:
  - Characters **BS** (backspace, code point U+0008), **HT** (horizontal tab, code point U+0009), **LF** (line feed, code point U+000A), **FF** (form feed, code point U+000C), **CR** (carriage return, code point U+000D), " (quotation mark, code point U+0022), and \ (backslash, code point U+005C) *MUST* be encoded using ECHAR.
  - Characters in the range from U+0000 to U+0007, **VT** (vertical tab, code point U+000B), characters in the range from U+000E to U+001F, **DEL** (delete, code point U+007F), and characters not matching the Char production from [XML11] *MUST* be represented by UCHAR using a lowercase \u with 4 HEXes.
  - All characters not required to be represented by ECHAR or UCHAR *MUST* be represented by their native [UNICODE] representation.
- The token EOL *MUST* be a single **LF** (line feed, code point U+000A).
- The final EOL *MUST* be provided.

## § B. URDNA2015

*This section is non-normative.*

RDF Dataset Canonicalization [CCG-RDC-FINAL] describes "Universal RDF Dataset Normalization Algorithm 2015" (URDNA2015), essentially the same algorithm as RDFC-1.0, and generally implementations implementing URDNA2015 should be compatible with this specification. The minor change is in the canonical n-quads form where some control characters were previously represented without escaping. The version of the algorithm defined in A. A Canonical form of N-Quads clarifies the representation of simple literals and the characters within STRING\_LITERAL\_QUOTE that are encoded using ECHAR.

## § C. URGNA2012

*This section is non-normative.*

A previous version of this algorithm has light deployment. For purposes of identification, the algorithm is called the "Universal RDF Graph Canonicalization Algorithm 2012" (**URGNA2012**), and differs from the stated algorithm in the following ways:

- In [4.6 Hash First Degree Quads](#), if any blank node was used in the [graph name](#) position in the [quad](#), then the value was serialized using the special [blank node identifier](#), [g](#), instead of [z](#).
- In [4.7 Hash Related Blank Node](#), value of the [predicate](#) was not delimited by [<](#) and [>](#); there were no delimiters.
- In [4.8 Hash N-Degree Quads](#), the *position* parameter passed to the [Hash Related Blank Node algorithm](#) was instead modeled as a *direction* parameter, where it could have the value [p](#), for property, when the related blank node was a [subject](#) and the value [r](#), for reverse or reference, when the related blank node was an [object](#). Since [URGNA2012](#) only canonicalized graphs, not datasets, there was no use of the [graph name](#) position.
- In [4.8 Hash N-Degree Quads](#), building the  $H_n$  was done as follows:
  - 1) For each *quad* in *quads*:
    - 1.1) If the *quad's subject* is a [blank node](#) that does not match *identifier*, set *hash* to the result of the [Hash Related Blank Node algorithm](#), passing the [blank node identifier](#) for [subject](#) as *related*, *quad*, *issuer*, and [p](#) as *position*.
    - 1.2) Otherwise, if *quad's object* is a [blank node](#) that does not match *identifier*, set *hash* to the result of the [Hash Related Blank Node algorithm](#), passing the [blank node identifier](#) for [object](#) as *related*, *quad*, *issuer*, and [r](#) as *position*.
    - 1.3) Otherwise, continue to the next quad.
    - 1.4) Add a mapping of *hash* to the [blank node identifier](#) for the component that matched ([subject](#) or [object](#)) to  $H_n$ , adding an entry as necessary.

## § D. Index

### § D.1 Terms defined by this specification

[blank node to quads map](#) §4.2

[canonical issuer](#) §4.2

[canonical n-quads form](#) §3.1

[canonicalization function](#) §3.1

[canonicalization state](#) §4.2

[canonicalized dataset](#) §3.1

<a href="#">definition</a>	§1.3	<a href="#">issued identifiers map</a>	§4.3
<a href="#">gossip path</a>	§3.1	<a href="#">mention</a>	§3.1
<a href="#">graph isomorphism</a>	§1.	<a href="#">mention set</a>	§3.1
<a href="#">hash</a>	§3.1	<a href="#">nquads</a>	§4.6.3
<a href="#">hash algorithm</a>	§3.1	<a href="#">partial canonicalization function</a>	§3.1
<a href="#">hash to blank nodes map</a>	§4.2	<a href="#">quad</a>	§3.1
<a href="#">identifier counter</a>	§4.3	<a href="#">RDFC-1.0</a>	§4.
<a href="#">identifier issuer</a>	§3.1	<a href="#">reference blank node identifier</a>	§4.6.3
<a href="#">identifier prefix</a>	§4.3	<a href="#">serialized canonical form</a>	§5.
<a href="#">input blank node identifier map</a>	§3.1	<a href="#">URDNA2015</a>	§B.
<a href="#">input dataset</a>	§3.1	<a href="#">URGNA2012</a>	§C.

## § D.2 Terms defined by reference

[BCP47] defines the following:

formatting conventions

[INFRA] defines the following:

boolean type

entry (for map)

key (for map)

list

map

string

[N-TRIPLES] defines the following:

Canonical N-Triples

[N-QUADS] defines the following:

C. N-Quads Internet Media Type, File Extension  
and Macintosh File Type

ECHAR

EOL

graphLabel

HEX

literal

STRING\_LITERAL\_QUOTE

UCHAR

[RDF11-CONCEPTS] defines the following:

- blank node identifiers
- blank nodes
- Blank Nodes
- default graph
- graph name
- IRIs
- isomorphic datasets
- language tags
- language-tagged strings
- lexical form
- literal
- literal term equality
- literal value
- object type
- predicate
- RDF datasets
- RDF graph
- RDF source
- RDF triple
- Section 3.3
- simple literals
- Skolem IRIs
- subject

[RDF11-MT] defines the following:

- RDF Collections

[VC-DATA-INTEGRITY] defines the following:

- Selective Disclosure

[XML11] defines the following:

- Char

[XPATH-FUNCTIONS] defines the following:

- Unicode code point order

## § E. Changes since the First Public Working Draft of 24 November 2022

*This section is non-normative.*

- The algorithm, and the examples, have been changed to systematically use the **xyz** format for blank node identifiers, instead of **\_:xyz**. See [Issue 46](#) for the discussion.
- [4.6 Hash First Degree Quads](#) was simplified to remove the **simple** flag, which was unused in existing implementations. The original design of the algorithm was to use the assigned canonical **blank node identifier**, if available, instead of **\_:a** or **\_:z**, similar to how it is used in the related hash algorithm, but this text never made it into the spec before implementations moved forward.

Therefore, the hashes never change, making the loop based on the **simple** flag that calls this algorithm unnecessary. See [Issue 23](#) for the discussion.

- Add definition for [canonical n-quads form](#). Eventually, this should be a citation from [N-Quads], when it is updated. [Canonical n-quads form](#) is used in [4.6 Hash First Degree Quads](#).
- Removed issue marker for [Issue 15](#) in [4.4 Canonicalization Algorithm](#), adding a note that [literal](#) components of [quads](#) are not normalized, and two literals with different syntactic representations remain distinct resources.
- Changed the way Blank Node identifiers are described (see [Issue 46](#)), generally omitting the leading `_:` which is a serialization artifact. This is still required in the algorithms, but the distinction between what is an identifier, and the serialization form is clarified.
- Changed the name of the algorithm from [URDNA2015](#) to [RDFC-1.0](#).
- Changed the term **normalized dataset** to [canonicalized dataset](#), which is composed of the [input dataset](#), [input blank node identifier map](#), and [issued identifiers map](#).

## § F. Changes since the Candidate Recommendation Snapshot of 31 October 2023

*This section is non-normative.*

- Clarified that detecting a poison dataset will result in an exception and early termination of the [Canonicalization Algorithm](#).

## § G. Acknowledgements

*This section is non-normative.*

The editors would like to thank Jeremy Carroll for his work on the graph canonicalization problem, Andy Seaborne and Gavin Carothers for providing valuable feedback and testing input for the algorithm defined in this specification, Sir Tim Berners-Lee for his thoughts on graph canonicalization over the years, Jesús Arias Fisteus for his work on a similar algorithm, and Aiden Hogan, whose publication [[Hogan-Canonical-RDF](#)] provided an important contemporary analysis of the canonicalization problem and served as an independent justification of the development of [RDFC-1.0](#).

The editors would also like to thank the chairs of the Working Group, Phil Archer and Markus Sabadello, and specific members of the Working Group whose active contributions were critical in completing this work: Pierre-Antoine Champin, Ivan Herman, David Lehn, Kazue Sako, Manu Sporny, and Ted Thibodeau Jr.

Members of the RDF Dataset Canonicalization and Hash Working Group included Ahamed Azeem, Ahmad Alabaid, Andy Seaborne, Benjamin Goering, Benjamin Young, Brent Zundel, Damien Graux, Dan Brickley, Dan Yamamoto, Daniel Pape, Dave Longley, David Lehn, Duy-Tung Luu, Gregg Kellogg, Ivan Herman, Jean-Yves Rossi, Jennifer Meier, Jesse Wright, Kazue Sako, Leonard Rosenthal, Mahmoud Alkhraishi, Manu Sporny, Markus Sabadello, Michael Prorock, Phil Archer, Pierre-Antoine Champin, Sebastian Crane, Ted Thibodeau Jr, Timothée Haudebourg, and Tobias Kuhn.

This specification is based on work done in the [W3C Credentials Community Group](#) published as [\[CCG-RDC-FINAL\]](#). Contributors to the Community Group Final Report include: Blake Regalia, Dave Longley, David Lehn, David Lozano Jarque, Gregg Kellogg, Manu Sporny, Markus Sabadello, Matt Collier, and Sebastian Schmittner.

Portions of the work on this specification have been funded by the European Union's StandICT.eu 2023 program under sub-grantee contract numbers No. 08/12 and 09/25. The content of this specification does not necessarily reflect the position or the policy of the European Union and no official endorsement should be inferred.

Portions of the work on this specification have also been funded by the U.S. Department of Homeland Security's Silicon Valley Innovation Program under contracts 70RSAT21T00000020 and 70RSAT23T00000006. The content of this specification does not necessarily reflect the position or the policy of the U.S. Government and no official endorsement should be inferred.

The Working Group acknowledges that the success of this specification is dependent on a long history of work performed over multiple decades in both academia and industry. We thank the individuals who iterated on the science which led to the completion of this specification. A partial list of these papers is found below, to the best of the Working Group's recollection. Omission from this list, whether intentional or unintentional, is not meant to imply that such an unlisted paper was not similarly important to the development of this work.

- Berners-Lee, T., Connolly, D.: [\*Delta: an ontology for the distribution of differences between RDF graphs\*](#) W3C Unofficial (2001). <https://www.w3.org/DesignIssues/Diff> [DesignIssues-Diff].
- Carroll, J.J.: [\*Signing RDF Graphs\*](#). In: Fensel, D., Sycara, K., and Mylopoulos, J. (eds.) The Semantic Web - ISWC 2003. pp. 369–384. Springer, Berlin, Heidelberg (2003).  
[https://doi.org/10.1007/978-3-540-39718-2\\_24](https://doi.org/10.1007/978-3-540-39718-2_24) [HPL-2003-142].

- Sayers, C., Karp, A.: Computing the digest of an RDF graph, Tech. Rep. HPL-2003-235 (R. 1), Hewlett Packard Laboratories (2004). <https://www.hpl.hp.com/techreports/2003/HPL-2003-235R1.html>.
- Carroll, J.J., Bizer, C., Hayes, P., Stickler, P.: Named graphs. Journal of Web Semantics. 3, 247–267 (2005). <https://doi.org/10.1016/j.websem.2005.09.001>.
- Tummarello, G., Morbidoni, C., Bachmann-Gmür, R., Erling, O.: RDFSync: Efficient Remote Synchronization of RDF Models. In: Aberer, K., Choi, K.-S., Noy, N., Allemang, D., Lee, K.-I., Nixon, L., Golbeck, J., Mika, P., Maynard, D., Mizoguchi, R., Schreiber, G., and Cudré-Mauroux, P. (eds.) The Semantic Web. pp. 537–551. Springer, Berlin, Heidelberg (2007). [https://doi.org/10.1007/978-3-540-76298-0\\_39](https://doi.org/10.1007/978-3-540-76298-0_39).
- Oren, E., Delbru, R., Catasta, M., Cyganiak, R., Stenzhorn, H., Tummarello, G.: Sindice.com: a document-oriented lookup index for open linked data. International Journal of Metadata, Semantics and Ontologies. 3, 37–52 (2008). <https://doi.org/10.1504/IJMSO.2008.021204>.
- Fisteus, J.A., Fernández García, N., Sánchez Fernández, L., Delgado Kloos, C.: Hashing and canonicalizing Notation 3 graphs. Journal of Computer and System Sciences. 76, 663–685 (2010). <https://doi.org/10.1016/j.jcss.2010.01.003>.
- Kasten, A., Scherp, A., Schauß, P.: *A Framework for Iterative Signing of Graph Data on the Web*. In: Presutti, V., d'Amato, C., Gandon, F., d'Aquin, M., Staab, S., and Tordai, A. (eds.) The Semantic Web: Trends and Challenges. pp. 146–160. Springer International Publishing, Cham (2014). [https://doi.org/10.1007/978-3-319-07443-6\\_11](https://doi.org/10.1007/978-3-319-07443-6_11) [eswc2014Kasten].
- Hogan, A.: *Canonical Forms for Isomorphic and Equivalent RDF Graphs: Algorithms for Leaning and Labelling Blank Nodes*. ACM Trans. Web. 11, 22:1-22:62 (2017). <https://doi.org/10.1145/3068333> [Hogan-Canonical-RDF].
- Arnold, R., Longley, D.: RDF Dataset Normalization. Report submitted to the W3C Credentials Community Group mailing list (2020). <https://lists.w3.org/Archives/Public/public-credentials/2021Mar/att-0220/RDFDatasetCanonicalization-2020-10-09.pdf>.

## § H. References

### § H.1 Normative references

#### [FIPS-180-4]

*FIPS PUB 180-4: Secure Hash Standard (SHS)*. U.S. Department of Commerce/National Institute of Standards and Technology. August 2015. National Standard. URL: <https://nvlpubs.nist.gov/nistpubs/FIPS/NIST.FIPS.180-4.pdf>

#### [INFRA]

*Infra Standard*. Anne van Kesteren; Domenic Denicola. WHATWG. Living Standard. URL: <https://infra.spec.whatwg.org/>

#### [N-Quads]

*RDF 1.1 N-Quads*. Gavin Carothers. W3C. 25 February 2014. W3C Recommendation. URL: <https://www.w3.org/TR/n-quads/>

#### [N-TRIPLES]

*RDF 1.1 N-Triples*. Gavin Carothers; Andy Seaborne. W3C. 25 February 2014. W3C Recommendation. URL: <https://www.w3.org/TR/n-triples/>

#### [RDF11-CONCEPTS]

*RDF 1.1 Concepts and Abstract Syntax*. Richard Cyganiak; David Wood; Markus Lanthaler. W3C. 25 February 2014. W3C Recommendation. URL: <https://www.w3.org/TR/rdf11-concepts/>

#### [RFC2119]

*Key words for use in RFCs to Indicate Requirement Levels*. S. Bradner. IETF. March 1997. Best Current Practice. URL: <https://www.rfc-editor.org/rfc/rfc2119>

#### [RFC3987]

*Internationalized Resource Identifiers (IRIs)*. M. Duerst; M. Suignard. IETF. January 2005. Proposed Standard. URL: <https://www.rfc-editor.org/rfc/rfc3987>

#### [RFC8174]

*Ambiguity of Uppercase vs Lowercase in RFC 2119 Key Words*. B. Leiba. IETF. May 2017. Best Current Practice. URL: <https://www.rfc-editor.org/rfc/rfc8174>

#### [Turtle]

*RDF 1.1 Turtle*. Eric Prud'hommeaux; Gavin Carothers. W3C. 25 February 2014. W3C Recommendation. URL: <https://www.w3.org/TR/turtle/>

**[UNICODE]**

*The Unicode Standard*. Unicode Consortium. URL: <https://www.unicode.org/versions/latest/>

**[XML11]**

*Extensible Markup Language (XML) 1.1 (Second Edition)*. Tim Bray; Jean Paoli; Michael Sperberg-McQueen; Eve Maler; François Yergeau; John Cowan et al. W3C. 16 August 2006. W3C Recommendation. URL: <https://www.w3.org/TR/xml11/>

**[XPATH-FUNCTIONS]**

*XQuery 1.0 and XPath 2.0 Functions and Operators (Second Edition)*. Ashok Malhotra; Jim Melton; Norman Walsh; Michael Kay. W3C. 14 December 2010. W3C Recommendation. URL: <https://www.w3.org/TR/xpath-functions/>

## § H.2 Informative references

**[BCP47]**

*Tags for Identifying Languages*. A. Phillips, Ed.; M. Davis, Ed.. IETF. September 2009. Best Current Practice. URL: <https://www.rfc-editor.org/rfc/rfc5646>

**[CCG-RDC-FINAL]**

*RDF Dataset Canonicalization*. Dave Longley. W3C. October 9, 2022. CG-FINAL. URL: <https://www.w3.org/community/reports/credentials/CG-FINAL-rdf-dataset-canonicalization-20221009/>

**[DesignIssues-Diff]**

*Delta: an ontology for the distribution of differences between RDF graphs*. Tim Berners-Lee. W3C. September 25, 2015. unofficial. URL: <https://www.w3.org/DesignIssues/Diff>

**[eswc2014Kasten]**

*A Framework for Iterative Signing of Graph Data on the Web*. Andreas Kasten; Ansgar Scherp; Peter Schauß . ISWC 2014. 2014. unofficial. URL: [https://doi.org/10.1007/978-3-319-07443-6\\_11](https://doi.org/10.1007/978-3-319-07443-6_11)

**[Hogan-Canonical-RDF]**

*Canonical Forms for Isomorphic and Equivalent RDF Graphs: Algorithms for Leaning and Labelling Blank Nodes*. Aiden Hogan. ACM. November 2017. ACM Trans. Web 11, 4, Article 22. URL: <https://aidanhogan.com/docs/rdf-canonicalisation.pdf>

**[HPL-2003-142]**

*Signing RDF Graphs*. Jeremy J. Carroll. HP Laboratories Bristol. July 23, 2003. unofficial. URL: <https://web.archive.org/web/20230129125726/https://www.hpl.hp.com/techreports/2003/HPL-2003-142.pdf>

**[RCH-EXPLAINER]**

*RDF Dataset Canonicalization and Hash Working Group — Explainer and Use Cases.* Phil Archer. W3C. 19 October 2023. W3C Working Group Note. URL: <https://www.w3.org/TR/rch-explainer/>

**[RDF11-MT]**

*RDF 1.1 Semantics.* Patrick Hayes; Peter Patel-Schneider. W3C. 25 February 2014. W3C Recommendation. URL: <https://www.w3.org/TR/rdf11-mt/>

**[VC-DATA-INTEGRITY]**

*Verifiable Credential Data Integrity 1.0.* Manu Sporny; Dave Longley; Greg Bernstein; Dmitri Zagidulin; Sebastian Crane. W3C. 28 April 2024. W3C Candidate Recommendation. URL: <https://www.w3.org/TR/vc-data-integrity/>

**[YAML]**

*YAML Ain't Markup Language (YAML™) Version 1.2.* Oren Ben-Kiki; Clark Evans; Ingy döt Net. 1 October 2009. URL: <http://yaml.org/spec/1.2/spec.html>

